

endo_Basic_Unet_mk1

May 11, 2025

```
[42]: import os
import numpy as np
from numpy.lib.stride_tricks import as_strided
import time
import matplotlib.pyplot as plt
from scipy.spatial.distance import directed_hausdorff

import torch
from torch.utils.data import DataLoader
from torch.utils.data import Dataset
from torch.utils.data import DataLoader, random_split
from torch.optim.lr_scheduler import StepLR

from pytorch_lightning import LightningDataModule
from pytorch_lightning import LightningModule
from pytorch_lightning import Trainer
from pytorch_lightning.callbacks import LearningRateMonitor, ModelCheckpoint
from pytorch_lightning.callbacks import EarlyStopping
from pytorch_lightning.loggers import TensorBoardLogger

from sklearn.model_selection import train_test_split

from monai.networks.nets import BasicUNet
from monai.losses import DiceCELoss
from monai.metrics import DiceMetric, MeanIoU, HausdorffDistanceMetric
from monai.transforms import (
    AsDiscreted,
    Compose,
    Resized,
    EnsureChannelFirstd,
    LoadImaged,
    ScaleIntensityd,
    ToTensord,
    RandFlipd,
    RandZoomd,
    ToTensord,
    AsDiscreted,
```

```

        CenterSpatialCropd
    )

```

```

[43]: # Custom dataset class for pytorch compatibility
# https://pytorch.org/tutorials/beginner/data_loading_tutorial.html
class EndoVis2017Dataset(Dataset):
    def __init__(self, label_subdir=None):
        self.data = []

        if label_subdir is None:
            raise ValueError("You must specify a `label_subdir` for ground_
↳truth masks (e.g., 'instrument_seg_composite').")

        self.root_dir = "C:/Users/dsumm/OneDrive/Documents/UMD ENPM Robotics_
↳Files/BIOE658B (Intro to Medical Image Analysis)/Project/dataset/train/"
        self.label_subdir = label_subdir

        # Recursively walk through directory to find left frame image paths and_
↳GT image paths
        for subdir, dirs, files in os.walk(self.root_dir):
            if 'left_frames' in subdir:
                #print("Hit!")
                for file in sorted(files):
                    if file.endswith(('.png', '.jpg', '.jpeg')):
                        img_path = os.path.join(subdir, file)
                        #print(img_path)

                        gt_root = subdir.replace('left_frames', 'ground_truth')
                        mask_path = os.path.join(gt_root, self.label_subdir,
↳file)

                        if os.path.exists(mask_path):
                            #print("Hit!")
                            self.data.append({"image": img_path, "label":
↳mask_path})    # Dictionary for MONAI compatability

        transforms_list = [
            LoadImaged(keys=["image", "label"]),    # Loads_
↳image data and metadata from file path dictionaries
            EnsureChannelFirstd(keys=["image", "label"]),    #
↳Adjust or add the channel dimension of input data to ensure channel_first_
↳shape

            # Images are of nominal size 1280x1024 --> resizing for memory_
↳efficiency

```

```

        CenterSpatialCropd(keys=["image", "label"], roi_size=(1024, 1280)),  

        # Cropping background padding from images  

        Resized(keys=["image", "label"], spatial_size=(256, 320)),  

        # Imported images are of various sizes: standardize to 320,256  

        # Apply data augmentation techniques  

        RandFlipd(keys=["image", "label"], prob=0.3, spatial_axis=1),  

        # Horizontal axis flip imposed w/ 30% prob  

        #RandRotate90d(keys=["image", "label"], prob=0.3, max_k=3),  

        # Random 90° rotation imposed w/ 30% prob  

        RandZoomd(keys=["image", "label"], prob=0.3, min_zoom=0.75,  

        max_zoom=1.25), # Zoom range (+/-25%) imposed w/ 30% prob  

        #RandAdjustContrastd(keys=["image"], prob=0.3, gamma=(0.75, 1.25)),  

        # Contrast variation (+/-25%) imposed w/ 30% prob  

        ScaleIntensityd(keys=["image"]), # Scale  

        # the intensity of input image to the value range 0-1  

        ToTensord(keys=["image", "label"]), #  

        # Ensure data is of tensor type for pytorch usage  

    ]  

    # Additional conditional transforms based on label_subdir  

    if label_subdir == "binary_composite":  

        transforms_list.append(AsDiscreted(keys=["label"], threshold=0.5))  

        # Binary threshold for binary seg  

    elif label_subdir == "part_seg_composite":  

        transforms_list.append(AsDiscreted(keys=["label"], to_onehot=5))  

        # 5 individual class labels for instrument independent part seg  

    elif label_subdir == "instrument_seg_composite":  

        transforms_list.append(AsDiscreted(keys=["label"], to_onehot=8))  

        # 8 individual class labels for part independent instrument seg  

    elif label_subdir == "instrument_part_seg_composite":  

        transforms_list.append(AsDiscreted(keys=["label"], to_onehot=21))  

        # 26 individual class labels for instrument & part seg  

    # Imposing MONAI transforms  

    # https://docs.monai.io/en/stable/transforms.html  

    self.transform = Compose(transforms_list)  

    def __len__(self):  

        # Returns number of imported samples  

        length = len(self.data)  

        return length  

    def __getitem__(self, idx):  

        # Return transformed sample from the dataset as dictated by the index

```

```

sample = self.data[idx]
return self.transform(sample)

```

```

[44]: class MONAIDataLoader(LightningDataModule):
    def __init__(self, dataset=None, batch_size: int = None, img_size: int = 256,
    ↪None, dimensions:int = None):
        super().__init__()
        if dataset is None:
            raise ValueError("No dataset given!")
        else:
            self.dataset = dataset

            self.train, self.val = random_split(self.dataset, [int(len(self.
    ↪dataset) * 0.8), len(self.dataset) - int(len(self.dataset) * 0.8)])
            self.batch_size = batch_size
            #self.num_workers = 2
            self.pin_memory = True
            #self.persistent_workers = True
            print(f"Train dataset size: {len(self.train)}")
            print(f"Validation dataset size: {len(self.val)}")

    def setup(self, stage=None):
        # required by PyTorch Lightning
        pass

    def train_dataloader(self):
        return DataLoader(self.train, batch_size=self.batch_size,
    ↪pin_memory=self.pin_memory)

    def val_dataloader(self):
        return DataLoader(self.val, batch_size=self.batch_size, pin_memory=self.
    ↪pin_memory)

    #def predict_dataloader(self):
    #    return DataLoader(self.test, batch_size=self.batch_size,
    ↪num_workers=16)

```

```

[45]: class basic_UNet_Train(LightningModule):
    def __init__(self, img_size=(1, 3, 256, 320), batch_size=1, lr=0.001,
    ↪num_classes=1):
        super().__init__()

        self.save_hyperparameters()
        self.num_classes = num_classes
        print("num_classes", self.num_classes, num_classes, self.hparams.
    ↪num_classes)
        self.example_input_array = [torch.zeros(self.hparams.img_size)]

```

```

        self.dice_metric = DiceMetric(include_background=True,
↪reduction="mean", ignore_empty=True)
        self.iou_metric = MeanIoU(include_background=True, reduction="mean",
↪ignore_empty=True)

        # Metric tracking
        self.dice_scores = []
        self.iou_scores = []

        # Defining MONAI Unet model paramters
        self.model = BasicUNet(spatial_dims=2,           # 2D image so spatial dims
↪= 2
                                in_channels=3,           # RGB input ultrasound image
                                out_channels=num_classes, # Binary
↪segmentation mask output image
                                features=(32, 64, 128, 256, 512, 32), #
↪standard Unet feature sizes (32, 32, 64, 128, 256, 32)
                                dropout=0.1)           # Dropout prob 10%

        # Using combined DICE and CE loss as loss function
        # Conditional loss function based on the number of classes
        if num_classes == 1:
            self.DICE_CE_Loss = DiceCELoss(
                include_background=False, # Exclude background class
                sigmoid=True, # Use softmax for multiclass segmentation
                softmax=False, # Apply softmax for multiclass
                lambda_dice=1.0, # Adjust the weight for Dice loss
                lambda_ce=1.0, # Adjust the weight for Cross-Entropy loss
                reduction='mean' # Use mean reduction
            )
        else:
            self.DICE_CE_Loss = DiceCELoss(
                include_background=False, # Exclude background class
                sigmoid=False, # Use softmax for multiclass segmentation
                softmax=True, # Apply softmax for multiclass
                lambda_dice=1.0, # Adjust the weight for Dice loss
                lambda_ce=1.0, # Adjust the weight for Cross-Entropy loss
                reduction='mean' # Use mean reduction
            )

        # Tracking losses for matplotlib
        self.train_losses = []
        self.val_losses = []

        # For storing images for the last epoch
        self.last_image = []

```

```

        self.last_pred = []
        self.last_mask = []
        self.logged_epochs = []

# Passes model inputs through U-net to get output predictions
    def forward(self, inputs):
        outputs = self.model(inputs)
        return outputs

    def training_step(self, batch, batch_idx):
# Gets labels for input and corresponding ground truth
        inputs, gt_input = self._prepare_batch(batch)

# Call forward pass
        outputs = self.forward(inputs)

# Compute DICE & CE loss based on current params
        loss = self.DICE_CE_Loss(outputs, gt_input)

# Log DICE loss with PyTorch Lightning logger
        self.log(f"Train_Dice_CE_loss", loss, on_epoch=True, prog_bar=True)

# Append train loss at the end of each epoch
        if batch_idx == len(batch) - 1:
            self.train_losses.append(loss.item())

        return loss

    def validation_step(self, batch, batch_idx):

# Gets labels for input and corresponding ground truth
        inputs, gt_input = self._prepare_batch(batch)
        outputs = self.forward(inputs)
        loss = self.DICE_CE_Loss(outputs, gt_input)
        self.log("val_loss", loss, on_step=False, on_epoch=True, prog_bar=True)

        if self.hparams.num_classes == 1:
            probs = torch.sigmoid(outputs)
            preds = (probs > 0.5).float()
# Ensure ground truth is binary (i.e., 0 or 1)
            gt_input = (gt_input > 0.5).float() # Threshold the ground truth

↪if needed

            intersection = (preds * gt_input).sum()
            union = preds.sum() + gt_input.sum()
            bin_dice_score = 2.0 * intersection / (union + 1e-8) # Avoid

↪division by zero

```

```

        # IoU score calculation for binary segmentation
        bin_iou_score = intersection / (union - intersection + 1e-8) #
        ↪ Avoid division by zero

        self.log("val_dice", bin_dice_score, on_step=False, on_epoch=True,
        ↪ prog_bar=True)
        self.log("val_iou", bin_iou_score, on_step=False, on_epoch=True,
        ↪ prog_bar=True)

    else:
        probs = torch.softmax(outputs, dim=1)
        preds = torch.nn.functional.one_hot(torch.argmax(probs, dim=1),
        ↪ num_classes=self.num_classes)
        preds = preds.permute(0, 3, 1, 2).float() # Shape: [B, C, H, W]

        self.dice_metric(y_pred=preds, y=gt_input)
        self.iou_metric(y_pred=preds, y=gt_input)

    if self.trainer.sanity_checking:
        return # skip logging during sanity check

    # Append validation loss at the end of each epoch
    if batch_idx == len(batch) - 1:
        self.val_losses.append(loss.item())

    # For binary segmentation: apply sigmoid and threshold
    if self.hparams.num_classes == 1:
        outputs = torch.sigmoid(outputs)
        outputs = (outputs > 0.5).float() # Convert probabilities to
        ↪ binary mask

        self.dice_scores.append(bin_dice_score)
        self.iou_scores.append(bin_iou_score)

    # For multiclass segmentation: apply softmax
    else:
        outputs = torch.softmax(outputs, dim=1) # Apply softmax for
        ↪ multi-class outputs
        dice = self.dice_metric.aggregate()[0].item()
        #print("Dice", dice)
        iou = self.iou_metric.aggregate()[0].item()
        #print("IOU", iou)
        self.dice_metric.reset()
        self.iou_metric.reset()
        self.dice_scores.append(dice)
        self.iou_scores.append(iou)

```

```

        self.log("val_dice", dice, on_step=False, on_epoch=True,
↪prog_bar=True)
        self.log("val_iou", iou, on_step=False, on_epoch=True,
↪prog_bar=True)

        # Normalize and convert tensor to 3 channels (RGB) for visualization
        def process(last):
            # Detach from cpu to not interrupt training
            # https://stackoverflow.com/questions/63582590/
↪why-do-we-call-detach-before-calling-numpy-on-a-pytorch-tensor
            last = last[0].detach().cpu()

            # Min max normalization
            # https://www.codecademy.com/article/normalization
            last = (last - last.min()) / (last.max() - last.min() + 1e-8)

            # If grayscale, reshape last image to RGB for display by
↪replicating gray value twice
            # https://discuss.pytorch.org/t/convert-grayscale-images-to-rgb/
↪113422
            return last.repeat(3, 1, 1) if last.shape[0] == 1 else last

        current_epoch = self.current_epoch
        total_epochs = self.trainer.max_epochs
        print("TE", total_epochs)

        if current_epoch == 0 or current_epoch == total_epochs - 1 or
↪current_epoch == total_epochs // 2:
            self.last_image.append(process(inputs))
            self.last_pred.append(process(outputs))
            self.last_mask.append(process(gt_input))
            self.logged_epochs.append(current_epoch)
            print(f"Logged image from epoch {current_epoch}")

        return loss

    #def predict_step(self, batch, batch_idx, dataloader_idx=0):
    #    return self(batch['image'])

    def configure_optimizers(self):
        #set optimizer
        optimizer = torch.optim.AdamW(self.parameters(), lr=self.hparams.lr,
↪weight_decay=1e-4)
        scheduler = StepLR(optimizer, step_size=5, gamma=0.5) # halve LR every
↪5 epochs
        return {

```



```

        'optimizer': optimizer,
        'lr_scheduler': {
            'scheduler': scheduler,
            'interval': 'epoch',
            'frequency': 1
        }
    }

def _prepare_batch(self, batch):
    return batch['image'], batch['label']

# Plot training and val losses when needed
def plot_losses(self):
    min_len = min(len(self.train_losses), len(self.val_losses))
    epochs = range(1, min_len + 1)

    # Plotting training vs validation loss
    plt.figure(figsize=(10, 6))
    plt.plot(epochs, self.train_losses[:len(epochs)], label="Training_
↪Loss", color='blue')
    plt.plot(epochs, self.val_losses[:len(epochs)], label="Validation_
↪Loss", color='orange')
    plt.title("Training vs Validation Loss")
    plt.xlabel("Epochs")
    plt.ylabel("Loss")
    plt.legend()
    plt.show()

def plot_metrics(self):
    epochs = range(1, len(self.dice_scores) + 1)

    # Convert to CPU floats if necessary
    dice = [d.cpu().item() if torch.is_tensor(d) else d for d in self.
↪dice_scores]
    iou = [i.cpu().item() if torch.is_tensor(i) else i for i in self.
↪iou_scores]

    plt.figure(figsize=(10, 6))
    plt.plot(epochs, dice, label='Dice Coefficient')
    plt.plot(epochs, iou, label='IoU')
    plt.xlabel("Epochs")
    plt.ylabel("Score")
    plt.title("Validation Metrics Over Time")
    plt.legend()
    plt.show()

def plot_result_by_epoch(self):

```

```

total_epochs = len(self.last_image)

if total_epochs < 5:
    print(f"Only {total_epochs} epochs recorded, plotting all.")
    selected_epochs = list(range(total_epochs))
else:
    print(f"{total_epochs} epochs recorded, bug in code.")

for epoch_idx in selected_epochs:
    epoch_num = self.logged_epochs[epoch_idx] if hasattr(self,
↪ "logged_epochs") else epoch_idx
    img = self.last_image[epoch_idx]
    pred = self.last_pred[epoch_idx]
    mask = self.last_mask[epoch_idx]

    fig, ax = plt.subplots(1, 3, figsize=(12, 4))

    ax[0].imshow(np.transpose(img.numpy(), (1, 2, 0)))
    ax[0].set_title(f"Epoch {epoch_num} - Image")
    ax[0].axis("off")

    if self.hparams.num_classes == 1:
        ax[1].imshow(np.transpose(pred.numpy(), (1, 2, 0)))
        ax[1].set_title(f"Epoch {epoch_num} - Prediction")
        ax[1].axis("off")

        ax[2].imshow(np.transpose(mask.numpy(), (1, 2, 0)))
        ax[2].set_title(f"Epoch {epoch_num} - Ground Truth")
        ax[2].axis("off")
    else:
        # Define the colormap and normalization
        num_classes = self.hparams.num_classes
        cmap = plt.get_cmap('viridis', num_classes)
        bounds = np.arange(num_classes + 1) - 0.5
        norm = plt.matplotlib.colors.BoundaryNorm(bounds, cmap.N)

        # Convert one-hot encoded predictions and masks to
↪ single-channel class labels
        pred_mask = torch.argmax(pred, dim=0).cpu().numpy()
        true_mask = torch.argmax(mask, dim=0).cpu().numpy()

        # Apply consistent colormap and normalization
        im1 = ax[1].imshow(pred_mask, cmap=cmap, norm=norm)
        ax[1].set_title(f"Epoch {epoch_num} - Prediction")
        ax[1].axis("off")

        im2 = ax[2].imshow(true_mask, cmap=cmap, norm=norm)

```

```

ax[2].set_title(f"Epoch {epoch_num} - Ground Truth")
ax[2].axis("off")

im_for_cbar = im1 # just need one mappable

# Adjust layout to leave space at the bottom
fig.subplots_adjust(bottom=0.25) # tweak this if labels get cut
↳off

# Add a new axis below the plots for the colorbar
cbar_ax = fig.add_axes([0.1, 0.1, 0.8, 0.10]) # [left, bottom,
↳width, height]

cbar = fig.colorbar(im_for_cbar, cax=cbar_ax,
↳orientation='horizontal', ticks=np.arange(num_classes))

# Add colorbar below the plots
# cbar = fig.colorbar(im1, ax=ax.ravel().tolist(),
↳orientation='horizontal',
# ticks=np.arange(num_classes), pad=0.15, fraction=0.05)

# Set class labels
if num_classes == 5:
    cbar.ax.set_xticklabels(['Background', 'Shaft', 'Wrist',
↳'Claspers', 'Probe'])
elif num_classes == 8:
    cbar.ax.set_xticklabels(['Background', 'Bipolar Forceps',
↳'Prograsp Forceps', 'Large Needle Driver',
'Vessel Sealer', 'Grasping
↳Retractor', 'Monopolar Curved Scissors', 'Other'])

plt.setp(cbar.ax.get_xticklabels(), rotation=30,
↳ha="right", rotation_mode="anchor")
elif num_classes == 21:
    cbar.ax.set_xticklabels([
        "Background",
        "Bipolar Forceps Shaft", "Bipolar Forceps Wrist",
↳"Bipolar Forceps Claspers",
        "Prograsp Forceps Shaft", "Prograsp Forceps Wrist",
↳"Prograsp Forceps Claspers",
        "Large Needle Driver Shaft", "Large Needle Driver
↳Wrist", "Large Needle Driver Claspers",
        "Vessel Sealer Shaft", "Vessel Sealer Wrist", "Vessel
↳Sealer Claspers",
        "Grasping Retractor Shaft", "Grasping Retractor Wrist",
↳"Grasping Retractor Claspers",

```

```

        "Monopolar Curved Scissors Shaft", "Monopolar Curved_
↪Scissors Wrist", "Monopolar Curved Scissors Claspers",
        "Other Probe", "Other Probe"
    ])
    plt.setp(cbar.ax.get_xticklabels(), rotation=45,
↪ha="right", rotation_mode="anchor")

    cbar.set_label('Class ID')

plt.show()

```

```

[46]: # Generate datasets, loaders, and models for basic UNet
binary_endo_images = EndoVis2017Dataset(label_subdir='binary_composite')
part_seg_endo_images = EndoVis2017Dataset(label_subdir='part_seg_composite')
instr_seg_endo_images =
↪EndoVis2017Dataset(label_subdir='instrument_seg_composite')
part_instr_seg_endo_images =
↪EndoVis2017Dataset(label_subdir='instrument_part_seg_composite')

binary_endo_data = MONAIDataLoader(dataset=binary_endo_images, batch_size=10)
↪# batch size should be divisible, ie. 50 images and bs 20 wont work
part_seg_endo_data = MONAIDataLoader(dataset=part_seg_endo_images,
↪batch_size=10)
instr_seg_endo_data = MONAIDataLoader(dataset=instr_seg_endo_images,
↪batch_size=10)
part_instr_seg_endo_data = MONAIDataLoader(dataset=part_instr_seg_endo_images,
↪batch_size=10)

binary_basic_UNet_model = basic_UNet_Train(num_classes=1)
part_seg_basic_UNet_model = basic_UNet_Train(num_classes=5)
instr_seg_basic_UNet_model = basic_UNet_Train(num_classes=8)
part_instr_seg_basic_UNet_model = basic_UNet_Train(num_classes=21)

```

```

Train dataset size: 1440
Validation dataset size: 360
Train dataset size: 1440
Validation dataset size: 360
Train dataset size: 1440
Validation dataset size: 360
Train dataset size: 1440
Validation dataset size: 360
num_classes 1 1 1
BasicUNet features: (32, 64, 128, 256, 512, 32).
num_classes 5 5 5
BasicUNet features: (32, 64, 128, 256, 512, 32).
num_classes 8 8 8
BasicUNet features: (32, 64, 128, 256, 512, 32).

```

```
num_classes 21 21 21
BasicUNet features: (32, 64, 128, 256, 512, 32).
```

```
[47]: if __name__ == "__main__":

    logger = TensorBoardLogger("tb_logs", name="binary_seg")

    early_stop_callback = EarlyStopping(
        monitor="Train_Dice_CE_loss",          # metric name from self.log
        mode="min",                            # because lower loss is better
        patience=5,                            # epochs to wait before stopping
        verbose=True
    )

    checkpoint_callback = ModelCheckpoint(
        monitor="Train_Dice_CE_loss",
        mode="min",
        save_top_k=1,
        dirpath="checkpoints/",
        filename="best-part-seg-basic-unet",
    )

    trainer = Trainer(
        accelerator="gpu",
        max_epochs=15,
        #limit_train_batches=0.1, # or 0.1 to use 10%
        logger=logger,
        callbacks=[early_stop_callback, checkpoint_callback],
    )

    start_train = time.time()
    trainer.fit(
        model=binary_basic_UNet_model,
        datamodule=binary_endo_data
    )
    end_train = time.time()
    print(f"Training time: {(end_train - start_train)/60:.2f} minutes")

    # Plot the overlaid training and val loss curves per epoch
    binary_basic_UNet_model.plot_losses()

    # Plot the IOU and DSC curves per epoch
    binary_basic_UNet_model.plot_metrics()

    # Plot images from last epoch
    binary_basic_UNet_model.plot_result_by_epoch()
```

```
GPU available: True (cuda), used: True
```

TPU available: False, using: 0 TPU cores
 HPU available: False, using: 0 HPUs
 Checkpoint directory C:\Users\dsumm\OneDrive\Documents\UMD ENPM Robotics
 Files\BIOE658B (Intro to Medical Image Analysis)\Project\code\checkpoints exists
 and is not empty.
 LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]

	Name	Type	Params	Mode	In sizes	Out sizes
0	model	BasicUNet	7.8 M	train	[1, 3, 256, 320]	[1, 1, 256, 320]
1	DICE_CE_Loss	DiceCELoss	0	train	?	?

7.8 M Trainable params
 0 Non-trainable params
 7.8 M Total params
 31.134 Total estimated model params size (MB)
 143 Modules in train mode
 0 Modules in eval mode

Sanity Checking: | 0/? [00:00<?, ?it/s]

The 'val_dataloader' does not have many workers which may be a bottleneck.
 Consider increasing the value of the `num_workers` argument` to `num_workers=31`
 in the `DataLoader` to improve performance.

Sanity Checking DataLoader 0: 50%| 1/2 [00:00<00:00, 13.15it/s]

single channel prediction, `include_background=False` ignored.

The 'train_dataloader' does not have many workers which may be a bottleneck.
 Consider increasing the value of the `num_workers` argument` to `num_workers=31`
 in the `DataLoader` to improve performance.

Epoch 0: 100%| 144/144 [01:43<00:00, 1.39it/s, v_num=5,
 Train_Dice_CE_loss_step=0.370]TE 15
 Logged image from epoch 0
 Epoch 0: 100%| 144/144 [02:00<00:00, 1.20it/s, v_num=5,
 Train_Dice_CE_loss_step=0.370, val_loss=0.458, val_dice=0.881, val_iou=0.789,
 Train_Dice_CE_loss_epoch=0.731]

Metric Train_Dice_CE_loss improved. New best score: 0.731

Epoch 1: 100%| 144/144 [01:50<00:00, 1.31it/s, v_num=5,
 Train_Dice_CE_loss_step=0.204, val_loss=0.458, val_dice=0.881, val_iou=0.789,
 Train_Dice_CE_loss_epoch=0.731]TE 15
 Epoch 1: 100%| 144/144 [02:06<00:00, 1.14it/s, v_num=5,

Train_Dice_CE_loss_step=0.204, val_loss=0.282, val_dice=0.898, val_iou=0.817,
Train_Dice_CE_loss_epoch=0.362]

Metric Train_Dice_CE_loss improved by 0.369 >= min_delta = 0.0. New best score:
0.362

Epoch 2: 100%| | 144/144 [01:31<00:00, 1.57it/s, v_num=5,
Train_Dice_CE_loss_step=0.184, val_loss=0.282, val_dice=0.898, val_iou=0.817,
Train_Dice_CE_loss_epoch=0.362]TE 15

Epoch 2: 100%| | 144/144 [01:46<00:00, 1.35it/s, v_num=5,
Train_Dice_CE_loss_step=0.184, val_loss=0.271, val_dice=0.885, val_iou=0.796,
Train_Dice_CE_loss_epoch=0.252]

Metric Train_Dice_CE_loss improved by 0.111 >= min_delta = 0.0. New best score:
0.252

Epoch 3: 100%| | 144/144 [01:30<00:00, 1.59it/s, v_num=5,
Train_Dice_CE_loss_step=0.117, val_loss=0.271, val_dice=0.885, val_iou=0.796,
Train_Dice_CE_loss_epoch=0.252]TE 15

Epoch 3: 100%| | 144/144 [01:46<00:00, 1.36it/s, v_num=5,
Train_Dice_CE_loss_step=0.117, val_loss=0.185, val_dice=0.924, val_iou=0.859,
Train_Dice_CE_loss_epoch=0.206]

Metric Train_Dice_CE_loss improved by 0.046 >= min_delta = 0.0. New best score:
0.206

Epoch 4: 100%| | 144/144 [01:31<00:00, 1.57it/s, v_num=5,
Train_Dice_CE_loss_step=0.128, val_loss=0.185, val_dice=0.924, val_iou=0.859,
Train_Dice_CE_loss_epoch=0.206]TE 15

Epoch 4: 100%| | 144/144 [01:47<00:00, 1.34it/s, v_num=5,
Train_Dice_CE_loss_step=0.128, val_loss=0.183, val_dice=0.922, val_iou=0.855,
Train_Dice_CE_loss_epoch=0.177]

Metric Train_Dice_CE_loss improved by 0.028 >= min_delta = 0.0. New best score:
0.177

Epoch 5: 100%| | 144/144 [01:35<00:00, 1.51it/s, v_num=5,
Train_Dice_CE_loss_step=0.0861, val_loss=0.183, val_dice=0.922, val_iou=0.855,
Train_Dice_CE_loss_epoch=0.177]TE 15

Epoch 5: 100%| | 144/144 [01:52<00:00, 1.28it/s, v_num=5,
Train_Dice_CE_loss_step=0.0861, val_loss=0.142, val_dice=0.944, val_iou=0.894,
Train_Dice_CE_loss_epoch=0.161]

Metric Train_Dice_CE_loss improved by 0.016 >= min_delta = 0.0. New best score:
0.161

Epoch 6: 100%| | 144/144 [01:42<00:00, 1.40it/s, v_num=5,
Train_Dice_CE_loss_step=0.079, val_loss=0.142, val_dice=0.944, val_iou=0.894,
Train_Dice_CE_loss_epoch=0.161] TE 15

Epoch 6: 100%| | 144/144 [01:59<00:00, 1.21it/s, v_num=5,
Train_Dice_CE_loss_step=0.079, val_loss=0.130, val_dice=0.948, val_iou=0.902,
Train_Dice_CE_loss_epoch=0.143]

Metric Train_Dice_CE_loss improved by 0.018 >= min_delta = 0.0. New best score: 0.143

Epoch 7: 100%| | 144/144 [01:41<00:00, 1.43it/s, v_num=5,
Train_Dice_CE_loss_step=0.0874, val_loss=0.130, val_dice=0.948, val_iou=0.902,
Train_Dice_CE_loss_epoch=0.143]TE 15

Logged image from epoch 7

Epoch 7: 100%| | 144/144 [01:57<00:00, 1.22it/s, v_num=5,
Train_Dice_CE_loss_step=0.0874, val_loss=0.123, val_dice=0.950, val_iou=0.905,
Train_Dice_CE_loss_epoch=0.141]

Metric Train_Dice_CE_loss improved by 0.002 >= min_delta = 0.0. New best score: 0.141

Epoch 8: 100%| | 144/144 [01:39<00:00, 1.45it/s, v_num=5,
Train_Dice_CE_loss_step=0.0855, val_loss=0.123, val_dice=0.950, val_iou=0.905,
Train_Dice_CE_loss_epoch=0.141]TE 15

Epoch 8: 100%| | 144/144 [01:54<00:00, 1.25it/s, v_num=5,
Train_Dice_CE_loss_step=0.0855, val_loss=0.115, val_dice=0.954, val_iou=0.912,
Train_Dice_CE_loss_epoch=0.129]

Metric Train_Dice_CE_loss improved by 0.012 >= min_delta = 0.0. New best score: 0.129

Epoch 9: 100%| | 144/144 [01:39<00:00, 1.45it/s, v_num=5,
Train_Dice_CE_loss_step=0.0776, val_loss=0.115, val_dice=0.954, val_iou=0.912,
Train_Dice_CE_loss_epoch=0.129]TE 15

Epoch 10: 100%| | 144/144 [01:38<00:00, 1.46it/s, v_num=5,
Train_Dice_CE_loss_step=0.0722, val_loss=0.109, val_dice=0.956, val_iou=0.916,
Train_Dice_CE_loss_epoch=0.129]TE 15

Epoch 10: 100%| | 144/144 [01:54<00:00, 1.26it/s, v_num=5,
Train_Dice_CE_loss_step=0.0722, val_loss=0.113, val_dice=0.954, val_iou=0.912,
Train_Dice_CE_loss_epoch=0.121]

Metric Train_Dice_CE_loss improved by 0.008 >= min_delta = 0.0. New best score: 0.121

Epoch 11: 100%| | 144/144 [01:34<00:00, 1.52it/s, v_num=5,
Train_Dice_CE_loss_step=0.0867, val_loss=0.113, val_dice=0.954, val_iou=0.912,
Train_Dice_CE_loss_epoch=0.121]TE 15

Epoch 11: 100%| | 144/144 [01:50<00:00, 1.31it/s, v_num=5,
Train_Dice_CE_loss_step=0.0867, val_loss=0.107, val_dice=0.956, val_iou=0.916,
Train_Dice_CE_loss_epoch=0.118]

Metric Train_Dice_CE_loss improved by 0.003 >= min_delta = 0.0. New best score: 0.118

Epoch 12: 100%| | 144/144 [01:36<00:00, 1.49it/s, v_num=5,
Train_Dice_CE_loss_step=0.0848, val_loss=0.107, val_dice=0.956, val_iou=0.916,
Train_Dice_CE_loss_epoch=0.118]TE 15

Epoch 12: 100%| | 144/144 [01:51<00:00, 1.29it/s, v_num=5,

Train_Dice_CE_loss_step=0.0848, val_loss=0.110, val_dice=0.956, val_iou=0.916,
Train_Dice_CE_loss_epoch=0.117]

Metric Train_Dice_CE_loss improved by 0.001 >= min_delta = 0.0. New best score:
0.117

Epoch 13: 100%| | 144/144 [01:39<00:00, 1.45it/s, v_num=5,
Train_Dice_CE_loss_step=0.0725, val_loss=0.110, val_dice=0.956, val_iou=0.916,
Train_Dice_CE_loss_epoch=0.117]TE 15

Epoch 13: 100%| | 144/144 [01:55<00:00, 1.24it/s, v_num=5,
Train_Dice_CE_loss_step=0.0725, val_loss=0.105, val_dice=0.958, val_iou=0.919,
Train_Dice_CE_loss_epoch=0.116]

Metric Train_Dice_CE_loss improved by 0.002 >= min_delta = 0.0. New best score:
0.116

Epoch 14: 100%| | 144/144 [01:39<00:00, 1.45it/s, v_num=5,
Train_Dice_CE_loss_step=0.0758, val_loss=0.105, val_dice=0.958, val_iou=0.919,
Train_Dice_CE_loss_epoch=0.116]TE 15

Logged image from epoch 14



Epoch 14: 100%| | 144/144 [01:55<00:00, 1.25it/s, v_num=5,
Train_Dice_CE_loss_step=0.0758, val_loss=0.102, val_dice=0.959, val_iou=0.921,
Train_Dice_CE_loss_epoch=0.111]

Metric Train_Dice_CE_loss improved by 0.004 >= min_delta = 0.0. New best score:
0.111

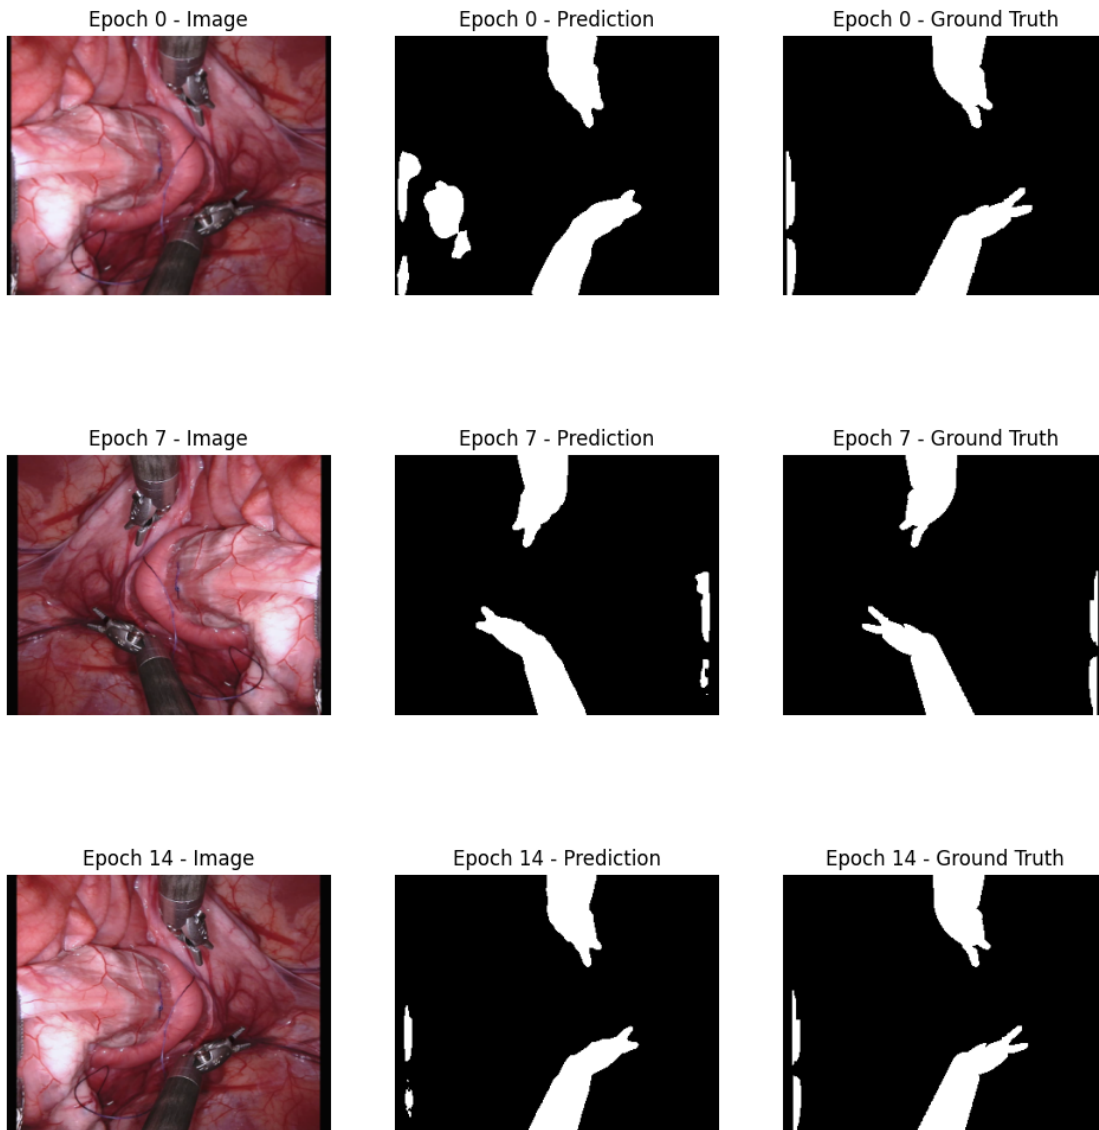
`Trainer.fit` stopped: `max_epochs=15` reached.

Epoch 14: 100%| | 144/144 [01:55<00:00, 1.24it/s, v_num=5,
Train_Dice_CE_loss_step=0.0758, val_loss=0.102, val_dice=0.959, val_iou=0.921,
Train_Dice_CE_loss_epoch=0.111]

Training time: 28.69 minutes



Only 3 epochs recorded, plotting all.



```
[48]: if __name__ == "__main__":

    logger = TensorBoardLogger("tb_logs", name="part_seg")

    early_stop_callback = EarlyStopping(
        monitor="Train_Dice_CE_loss",          # metric name from self.log
        mode="min",                            # because lower loss is better
        patience=5,                            # epochs to wait before stopping
        verbose=True
    )

    checkpoint_callback = ModelCheckpoint(
```

```

        monitor="Train_Dice_CE_loss",
        mode="min",
        save_top_k=1,
        dirpath="checkpoints/",
        filename="best-part-seg-basic-unet",
    )

    trainer = Trainer(
        accelerator="gpu",
        max_epochs=15,
        #limit_train_batches=0.1, # or 0.1 to use 10%
        logger=logger,
        callbacks=[early_stop_callback, checkpoint_callback],
    )

    start_train = time.time()
    trainer.fit(
        model=part_seg_basic_UNet_model,
        datamodule=part_seg_endo_data
    )
    end_train = time.time()
    print(f"Training time: {(end_train - start_train)/60:.2f} minutes")

    # Plot the overlaid training and val loss curves per epoch
    part_seg_basic_UNet_model.plot_losses()

    # Plot the IOU and DSC curves per epoch
    part_seg_basic_UNet_model.plot_metrics()

    # Plot images from last epoch
    part_seg_basic_UNet_model.plot_result_by_epoch()

```

GPU available: True (cuda), used: True
 TPU available: False, using: 0 TPU cores
 HPU available: False, using: 0 HPUs
 LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]

	Name	Type	Params	Mode	In sizes	Out sizes
0	model	BasicUNet	7.8 M	train	[1, 3, 256, 320]	[1, 5, 256, 320]
1	DICE_CE_Loss	DiceCELoss	0	train	?	?

7.8 M	Trainable params
0	Non-trainable params
7.8 M	Total params

```

31.134    Total estimated model params size (MB)
143      Modules in train mode
0        Modules in eval mode

Epoch 0: 100%|      | 144/144 [01:40<00:00,  1.43it/s, v_num=6,
Train_Dice_CE_loss_step=1.120]TE 15
Logged image from epoch 0
Epoch 0: 100%|      | 144/144 [02:00<00:00,  1.20it/s, v_num=6,
Train_Dice_CE_loss_step=1.120, val_loss=1.130, val_dice=0.291, val_iou=0.234,
Train_Dice_CE_loss_epoch=1.660]

Metric Train_Dice_CE_loss improved. New best score: 1.664

Epoch 1: 100%|      | 144/144 [01:39<00:00,  1.45it/s, v_num=6,
Train_Dice_CE_loss_step=0.857, val_loss=1.130, val_dice=0.291, val_iou=0.234,
Train_Dice_CE_loss_epoch=1.660]TE 15
Epoch 1: 100%|      | 144/144 [01:58<00:00,  1.21it/s, v_num=6,
Train_Dice_CE_loss_step=0.857, val_loss=0.913, val_dice=0.524, val_iou=0.449,
Train_Dice_CE_loss_epoch=0.980]

Metric Train_Dice_CE_loss improved by 0.684 >= min_delta = 0.0. New best score:
0.980

Epoch 2: 100%|      | 144/144 [01:36<00:00,  1.49it/s, v_num=6,
Train_Dice_CE_loss_step=0.690, val_loss=0.913, val_dice=0.524, val_iou=0.449,
Train_Dice_CE_loss_epoch=0.980]TE 15
Epoch 2: 100%|      | 144/144 [01:55<00:00,  1.24it/s, v_num=6,
Train_Dice_CE_loss_step=0.690, val_loss=0.804, val_dice=0.614, val_iou=0.523,
Train_Dice_CE_loss_epoch=0.828]

Metric Train_Dice_CE_loss improved by 0.152 >= min_delta = 0.0. New best score:
0.828

Epoch 3: 100%|      | 144/144 [01:37<00:00,  1.47it/s, v_num=6,
Train_Dice_CE_loss_step=0.681, val_loss=0.804, val_dice=0.614, val_iou=0.523,
Train_Dice_CE_loss_epoch=0.828]TE 15
Epoch 3: 100%|      | 144/144 [01:57<00:00,  1.22it/s, v_num=6,
Train_Dice_CE_loss_step=0.681, val_loss=0.709, val_dice=0.681, val_iou=0.583,
Train_Dice_CE_loss_epoch=0.749]

Metric Train_Dice_CE_loss improved by 0.079 >= min_delta = 0.0. New best score:
0.749

Epoch 4: 100%|      | 144/144 [01:41<00:00,  1.42it/s, v_num=6,
Train_Dice_CE_loss_step=0.630, val_loss=0.709, val_dice=0.681, val_iou=0.583,
Train_Dice_CE_loss_epoch=0.749]TE 15
Epoch 4: 100%|      | 144/144 [02:00<00:00,  1.19it/s, v_num=6,
Train_Dice_CE_loss_step=0.630, val_loss=0.649, val_dice=0.721, val_iou=0.626,
Train_Dice_CE_loss_epoch=0.672]

Metric Train_Dice_CE_loss improved by 0.077 >= min_delta = 0.0. New best score:
0.672

```

Epoch 5: 100%| | 144/144 [01:38<00:00, 1.46it/s, v_num=6,
Train_Dice_CE_loss_step=0.512, val_loss=0.649, val_dice=0.721, val_iou=0.626,
Train_Dice_CE_loss_epoch=0.672]TE 15
Epoch 5: 100%| | 144/144 [01:58<00:00, 1.22it/s, v_num=6,
Train_Dice_CE_loss_step=0.512, val_loss=0.596, val_dice=0.749, val_iou=0.657,
Train_Dice_CE_loss_epoch=0.615]

Metric Train_Dice_CE_loss improved by 0.057 >= min_delta = 0.0. New best score:
0.615

Epoch 6: 100%| | 144/144 [01:38<00:00, 1.46it/s, v_num=6,
Train_Dice_CE_loss_step=0.504, val_loss=0.596, val_dice=0.749, val_iou=0.657,
Train_Dice_CE_loss_epoch=0.615]TE 15
Epoch 6: 100%| | 144/144 [01:58<00:00, 1.22it/s, v_num=6,
Train_Dice_CE_loss_step=0.504, val_loss=0.610, val_dice=0.780, val_iou=0.691,
Train_Dice_CE_loss_epoch=0.582]

Metric Train_Dice_CE_loss improved by 0.033 >= min_delta = 0.0. New best score:
0.582

Epoch 7: 100%| | 144/144 [01:38<00:00, 1.46it/s, v_num=6,
Train_Dice_CE_loss_step=0.523, val_loss=0.610, val_dice=0.780, val_iou=0.691,
Train_Dice_CE_loss_epoch=0.582]TE 15
Logged image from epoch 7
Epoch 7: 100%| | 144/144 [01:58<00:00, 1.21it/s, v_num=6,
Train_Dice_CE_loss_step=0.523, val_loss=0.583, val_dice=0.769, val_iou=0.680,
Train_Dice_CE_loss_epoch=0.570]

Metric Train_Dice_CE_loss improved by 0.013 >= min_delta = 0.0. New best score:
0.570

Epoch 8: 100%| | 144/144 [01:39<00:00, 1.44it/s, v_num=6,
Train_Dice_CE_loss_step=0.458, val_loss=0.583, val_dice=0.769, val_iou=0.680,
Train_Dice_CE_loss_epoch=0.570]TE 15
Epoch 8: 100%| | 144/144 [01:59<00:00, 1.21it/s, v_num=6,
Train_Dice_CE_loss_step=0.458, val_loss=0.565, val_dice=0.782, val_iou=0.695,
Train_Dice_CE_loss_epoch=0.553]

Metric Train_Dice_CE_loss improved by 0.017 >= min_delta = 0.0. New best score:
0.553

Epoch 9: 100%| | 144/144 [01:39<00:00, 1.44it/s, v_num=6,
Train_Dice_CE_loss_step=0.435, val_loss=0.565, val_dice=0.782, val_iou=0.695,
Train_Dice_CE_loss_epoch=0.553]TE 15
Epoch 9: 100%| | 144/144 [01:59<00:00, 1.21it/s, v_num=6,
Train_Dice_CE_loss_step=0.435, val_loss=0.567, val_dice=0.792, val_iou=0.708,
Train_Dice_CE_loss_epoch=0.547]

Metric Train_Dice_CE_loss improved by 0.006 >= min_delta = 0.0. New best score:
0.547

Epoch 10: 100%| | 144/144 [01:39<00:00, 1.44it/s, v_num=6,
Train_Dice_CE_loss_step=0.422, val_loss=0.567, val_dice=0.792, val_iou=0.708,

Train_Dice_CE_loss_epoch=0.547]TE 15
Epoch 10: 100%| | 144/144 [01:59<00:00, 1.21it/s, v_num=6,
Train_Dice_CE_loss_step=0.422, val_loss=0.537, val_dice=0.794, val_iou=0.710,
Train_Dice_CE_loss_epoch=0.523]

Metric Train_Dice_CE_loss improved by 0.024 >= min_delta = 0.0. New best score:
0.523

Epoch 11: 100%| | 144/144 [01:40<00:00, 1.43it/s, v_num=6,
Train_Dice_CE_loss_step=0.443, val_loss=0.537, val_dice=0.794, val_iou=0.710,
Train_Dice_CE_loss_epoch=0.523]TE 15
Epoch 11: 100%| | 144/144 [02:00<00:00, 1.20it/s, v_num=6,
Train_Dice_CE_loss_step=0.443, val_loss=0.522, val_dice=0.804, val_iou=0.722,
Train_Dice_CE_loss_epoch=0.510]

Metric Train_Dice_CE_loss improved by 0.013 >= min_delta = 0.0. New best score:
0.510

Epoch 12: 100%| | 144/144 [01:40<00:00, 1.44it/s, v_num=6,
Train_Dice_CE_loss_step=0.467, val_loss=0.522, val_dice=0.804, val_iou=0.722,
Train_Dice_CE_loss_epoch=0.510]TE 15
Epoch 12: 100%| | 144/144 [01:59<00:00, 1.20it/s, v_num=6,
Train_Dice_CE_loss_step=0.467, val_loss=0.519, val_dice=0.811, val_iou=0.731,
Train_Dice_CE_loss_epoch=0.506]

Metric Train_Dice_CE_loss improved by 0.004 >= min_delta = 0.0. New best score:
0.506

Epoch 13: 100%| | 144/144 [01:40<00:00, 1.43it/s, v_num=6,
Train_Dice_CE_loss_step=0.420, val_loss=0.519, val_dice=0.811, val_iou=0.731,
Train_Dice_CE_loss_epoch=0.506]TE 15
Epoch 13: 100%| | 144/144 [01:59<00:00, 1.20it/s, v_num=6,
Train_Dice_CE_loss_step=0.420, val_loss=0.516, val_dice=0.813, val_iou=0.732,
Train_Dice_CE_loss_epoch=0.499]

Metric Train_Dice_CE_loss improved by 0.007 >= min_delta = 0.0. New best score:
0.499



Epoch 14: 100%| | 144/144 [01:39<00:00, 1.44it/s, v_num=6,
Train_Dice_CE_loss_step=0.431, val_loss=0.516, val_dice=0.813, val_iou=0.732,
Train_Dice_CE_loss_epoch=0.499]TE 15
Logged image from epoch 14
Epoch 14: 100%| | 144/144 [01:59<00:00, 1.20it/s, v_num=6,
Train_Dice_CE_loss_step=0.431, val_loss=0.517, val_dice=0.812, val_iou=0.733,
Train_Dice_CE_loss_epoch=0.495]

Metric Train_Dice_CE_loss improved by 0.003 >= min_delta = 0.0. New best score:
0.495

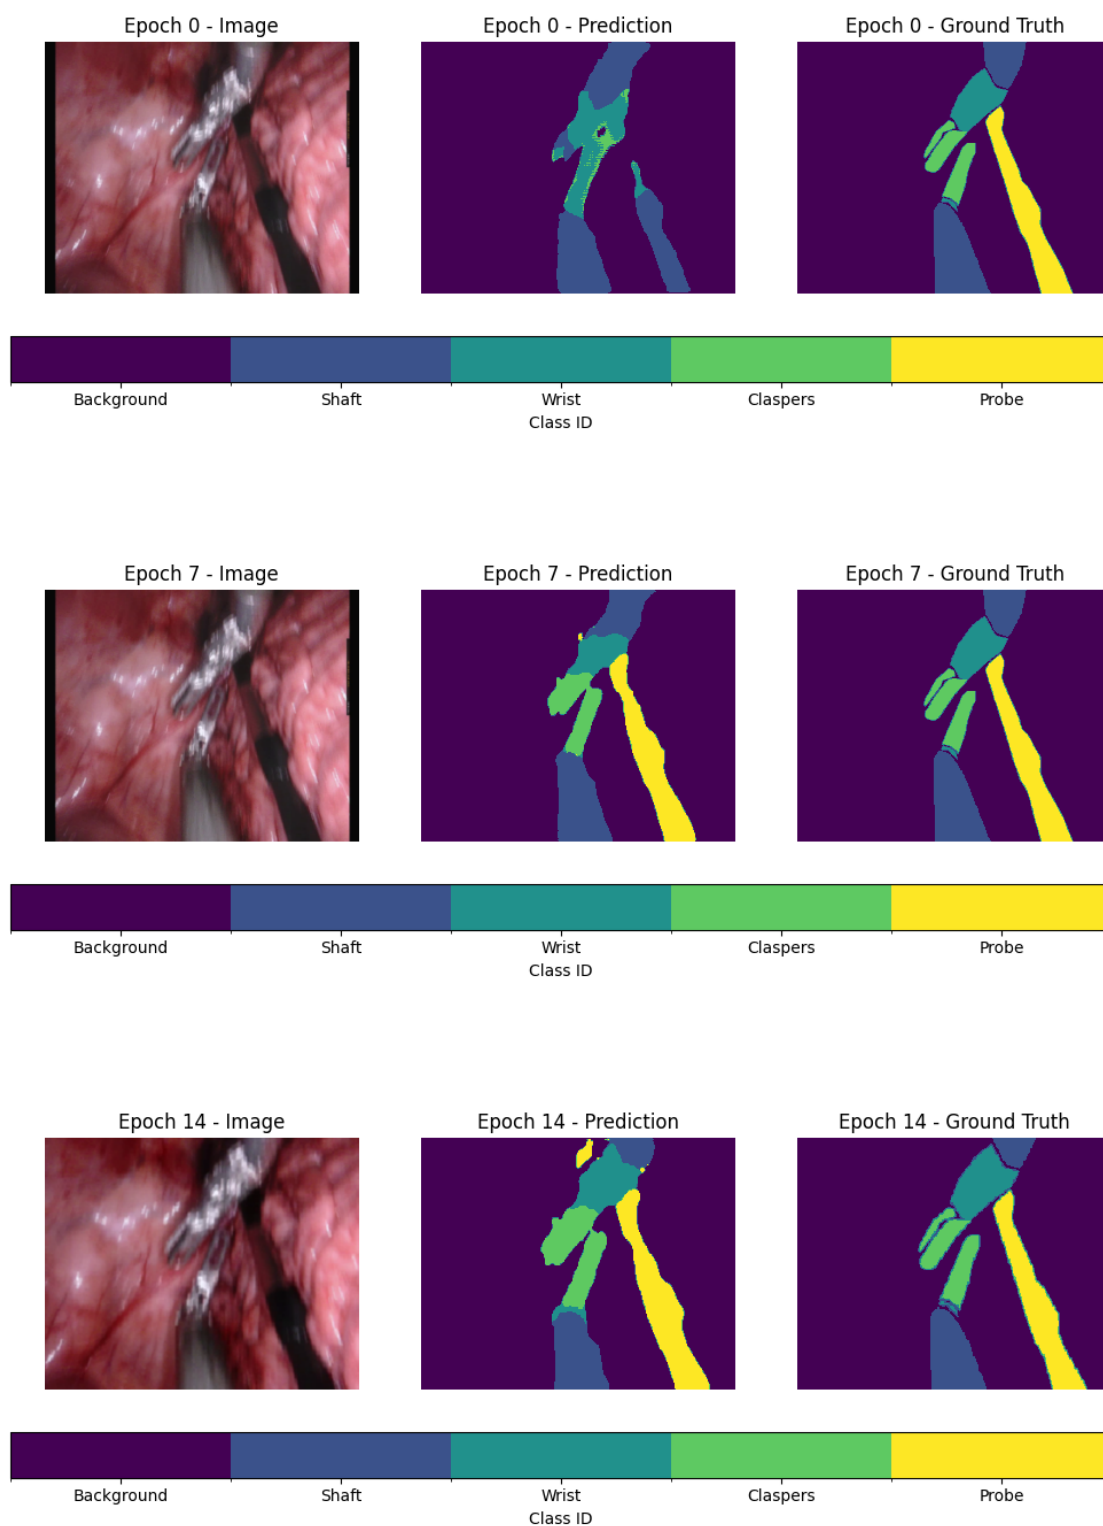
`Trainer.fit` stopped: `max_epochs=15` reached.

Epoch 14: 100%| | 144/144 [01:59<00:00, 1.20it/s, v_num=6,
Train_Dice_CE_loss_step=0.431, val_loss=0.517, val_dice=0.812, val_iou=0.733,

Train_Dice_CE_loss_epoch=0.495]
Training time: 29.88 minutes



Only 3 epochs recorded, plotting all.



```

[49]: if __name__ == "__main__":

    logger = TensorBoardLogger("tb_logs", name="instrument_seg")

    early_stop_callback = EarlyStopping(
        monitor="Train_Dice_CE_loss",          # metric name from self.log
        mode="min",                            # because lower loss is better
        patience=5,                            # epochs to wait before stopping
        verbose=True
    )

    checkpoint_callback = ModelCheckpoint(
        monitor="Train_Dice_CE_loss",
        mode="min",
        save_top_k=1,
        dirpath="checkpoints/",
        filename="best-instrument-seg-basic-unet",
    )

    trainer = Trainer(
        accelerator="gpu",
        max_epochs=20,
        #limit_train_batches=0.1, # or 0.1 to use 10%
        logger=logger,
        callbacks=[early_stop_callback, checkpoint_callback],
    )

    start_train = time.time()
    trainer.fit(
        model=instr_seg_basic_UNet_model,
        datamodule=instr_seg_endo_data
    )
    end_train = time.time()
    print(f"Training time: {(end_train - start_train)/60:.2f} minutes")

    # Plot the overlaid training and val loss curves per epoch
    instr_seg_basic_UNet_model.plot_losses()

    # Plot the IOU and DSC curves per epoch
    instr_seg_basic_UNet_model.plot_metrics()

    # Plot images from last epoch
    instr_seg_basic_UNet_model.plot_result_by_epoch()

```

GPU available: True (cuda), used: True
 TPU available: False, using: 0 TPU cores

HPU available: False, using: 0 HPUs
 Checkpoint directory C:\Users\dsumm\OneDrive\Documents\UMD ENPM Robotics
 Files\BIOE658B (Intro to Medical Image Analysis)\Project\code\checkpoints exists
 and is not empty.
 LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]

	Name	Type	Params	Mode	In sizes	Out sizes
0	model	BasicUNet	7.8 M	train	[1, 3, 256, 320]	[1, 8, 256, 320]
1	DICE_CE_Loss	DiceCELoss	0	train	?	?

7.8 M Trainable params
 0 Non-trainable params
 7.8 M Total params
 31.135 Total estimated model params size (MB)
 143 Modules in train mode
 0 Modules in eval mode

Sanity Checking: | 0/? [00:00<?, ?it/s]

The 'val_dataloader' does not have many workers which may be a bottleneck.
 Consider increasing the value of the `num_workers` argument` to `num_workers=31`
 in the `DataLoader` to improve performance.

The 'train_dataloader' does not have many workers which may be a bottleneck.
 Consider increasing the value of the `num_workers` argument` to `num_workers=31`
 in the `DataLoader` to improve performance.

Epoch 0: 100%| 144/144 [01:49<00:00, 1.32it/s, v_num=0,
 Train_Dice_CE_loss_step=1.360]TE 20
 Logged image from epoch 0
 Epoch 0: 100%| 144/144 [02:11<00:00, 1.10it/s, v_num=0,
 Train_Dice_CE_loss_step=1.360, val_loss=1.340, val_dice=0.217, val_iou=0.182,
 Train_Dice_CE_loss_epoch=1.700]

Metric Train_Dice_CE_loss improved. New best score: 1.695

Epoch 1: 100%| 144/144 [01:38<00:00, 1.46it/s, v_num=0,
 Train_Dice_CE_loss_step=1.280, val_loss=1.340, val_dice=0.217, val_iou=0.182,
 Train_Dice_CE_loss_epoch=1.700]TE 20
 Epoch 1: 100%| 144/144 [01:58<00:00, 1.21it/s, v_num=0,
 Train_Dice_CE_loss_step=1.280, val_loss=1.240, val_dice=0.270, val_iou=0.256,
 Train_Dice_CE_loss_epoch=1.300]

Metric Train_Dice_CE_loss improved by 0.395 >= min_delta = 0.0. New best score:
 1.300

Epoch 2: 100%| | 144/144 [01:39<00:00, 1.45it/s, v_num=0,
Train_Dice_CE_loss_step=1.210, val_loss=1.240, val_dice=0.270, val_iou=0.256,
Train_Dice_CE_loss_epoch=1.300]TE 20
Epoch 2: 100%| | 144/144 [01:58<00:00, 1.21it/s, v_num=0,
Train_Dice_CE_loss_step=1.210, val_loss=1.180, val_dice=0.317, val_iou=0.285,
Train_Dice_CE_loss_epoch=1.220]

Metric Train_Dice_CE_loss improved by 0.081 >= min_delta = 0.0. New best score:
1.219

Epoch 3: 100%| | 144/144 [01:39<00:00, 1.45it/s, v_num=0,
Train_Dice_CE_loss_step=1.170, val_loss=1.180, val_dice=0.317, val_iou=0.285,
Train_Dice_CE_loss_epoch=1.220]TE 20
Epoch 3: 100%| | 144/144 [01:58<00:00, 1.22it/s, v_num=0,
Train_Dice_CE_loss_step=1.170, val_loss=1.120, val_dice=0.362, val_iou=0.316,
Train_Dice_CE_loss_epoch=1.160]

Metric Train_Dice_CE_loss improved by 0.055 >= min_delta = 0.0. New best score:
1.164

Epoch 4: 100%| | 144/144 [01:34<00:00, 1.52it/s, v_num=0,
Train_Dice_CE_loss_step=1.080, val_loss=1.120, val_dice=0.362, val_iou=0.316,
Train_Dice_CE_loss_epoch=1.160]TE 20
Epoch 4: 100%| | 144/144 [01:53<00:00, 1.26it/s, v_num=0,
Train_Dice_CE_loss_step=1.080, val_loss=1.040, val_dice=0.392, val_iou=0.343,
Train_Dice_CE_loss_epoch=1.110]

Metric Train_Dice_CE_loss improved by 0.056 >= min_delta = 0.0. New best score:
1.108

Epoch 5: 100%| | 144/144 [01:35<00:00, 1.51it/s, v_num=0,
Train_Dice_CE_loss_step=1.020, val_loss=1.040, val_dice=0.392, val_iou=0.343,
Train_Dice_CE_loss_epoch=1.110]TE 20
Epoch 5: 100%| | 144/144 [01:54<00:00, 1.26it/s, v_num=0,
Train_Dice_CE_loss_step=1.020, val_loss=0.996, val_dice=0.420, val_iou=0.372,
Train_Dice_CE_loss_epoch=1.030]

Metric Train_Dice_CE_loss improved by 0.079 >= min_delta = 0.0. New best score:
1.029

Epoch 6: 100%| | 144/144 [01:36<00:00, 1.50it/s, v_num=0,
Train_Dice_CE_loss_step=1.000, val_loss=0.996, val_dice=0.420, val_iou=0.372,
Train_Dice_CE_loss_epoch=1.030]TE 20
Epoch 6: 100%| | 144/144 [01:55<00:00, 1.25it/s, v_num=0,
Train_Dice_CE_loss_step=1.000, val_loss=0.958, val_dice=0.441, val_iou=0.394,
Train_Dice_CE_loss_epoch=0.989]

Metric Train_Dice_CE_loss improved by 0.041 >= min_delta = 0.0. New best score:
0.989

Epoch 7: 100%| | 144/144 [01:36<00:00, 1.49it/s, v_num=0,
Train_Dice_CE_loss_step=1.010, val_loss=0.958, val_dice=0.441, val_iou=0.394,
Train_Dice_CE_loss_epoch=0.989]TE 20

Epoch 7: 100%| | 144/144 [01:55<00:00, 1.25it/s, v_num=0,
Train_Dice_CE_loss_step=1.010, val_loss=0.931, val_dice=0.453, val_iou=0.409,
Train_Dice_CE_loss_epoch=0.956]

Metric Train_Dice_CE_loss improved by 0.033 >= min_delta = 0.0. New best score:
0.956

Epoch 8: 100%| | 144/144 [01:37<00:00, 1.47it/s, v_num=0,
Train_Dice_CE_loss_step=1.010, val_loss=0.931, val_dice=0.453, val_iou=0.409,
Train_Dice_CE_loss_epoch=0.956]TE 20

Epoch 8: 100%| | 144/144 [01:57<00:00, 1.22it/s, v_num=0,
Train_Dice_CE_loss_step=1.010, val_loss=0.927, val_dice=0.460, val_iou=0.419,
Train_Dice_CE_loss_epoch=0.940]

Metric Train_Dice_CE_loss improved by 0.016 >= min_delta = 0.0. New best score:
0.940

Epoch 9: 100%| | 144/144 [01:36<00:00, 1.49it/s, v_num=0,
Train_Dice_CE_loss_step=0.950, val_loss=0.927, val_dice=0.460, val_iou=0.419,
Train_Dice_CE_loss_epoch=0.940]TE 20

Epoch 9: 100%| | 144/144 [01:55<00:00, 1.24it/s, v_num=0,
Train_Dice_CE_loss_step=0.950, val_loss=0.930, val_dice=0.466, val_iou=0.424,
Train_Dice_CE_loss_epoch=0.915]

Metric Train_Dice_CE_loss improved by 0.025 >= min_delta = 0.0. New best score:
0.915

Epoch 10: 100%| | 144/144 [01:37<00:00, 1.48it/s, v_num=0,
Train_Dice_CE_loss_step=0.936, val_loss=0.930, val_dice=0.466, val_iou=0.424,
Train_Dice_CE_loss_epoch=0.915]TE 20

Logged image from epoch 10

Epoch 10: 100%| | 144/144 [01:56<00:00, 1.23it/s, v_num=0,
Train_Dice_CE_loss_step=0.936, val_loss=0.893, val_dice=0.471, val_iou=0.427,
Train_Dice_CE_loss_epoch=0.890]

Metric Train_Dice_CE_loss improved by 0.025 >= min_delta = 0.0. New best score:
0.890

Epoch 11: 100%| | 144/144 [01:37<00:00, 1.47it/s, v_num=0,
Train_Dice_CE_loss_step=0.904, val_loss=0.893, val_dice=0.471, val_iou=0.427,
Train_Dice_CE_loss_epoch=0.890]TE 20

Epoch 11: 100%| | 144/144 [01:57<00:00, 1.23it/s, v_num=0,
Train_Dice_CE_loss_step=0.904, val_loss=0.883, val_dice=0.487, val_iou=0.442,
Train_Dice_CE_loss_epoch=0.877]

Metric Train_Dice_CE_loss improved by 0.013 >= min_delta = 0.0. New best score:
0.877

Epoch 12: 100%| | 144/144 [01:38<00:00, 1.46it/s, v_num=0,
Train_Dice_CE_loss_step=0.919, val_loss=0.883, val_dice=0.487, val_iou=0.442,
Train_Dice_CE_loss_epoch=0.877]TE 20

Epoch 12: 100%| | 144/144 [01:57<00:00, 1.22it/s, v_num=0,

Train_Dice_CE_loss_step=0.919, val_loss=0.883, val_dice=0.490, val_iou=0.446,
Train_Dice_CE_loss_epoch=0.872]

Metric Train_Dice_CE_loss improved by 0.005 >= min_delta = 0.0. New best score:
0.872

Epoch 13: 100%| | 144/144 [01:38<00:00, 1.46it/s, v_num=0,
Train_Dice_CE_loss_step=0.905, val_loss=0.883, val_dice=0.490, val_iou=0.446,
Train_Dice_CE_loss_epoch=0.872]TE 20

Epoch 13: 100%| | 144/144 [01:58<00:00, 1.22it/s, v_num=0,
Train_Dice_CE_loss_step=0.905, val_loss=0.871, val_dice=0.491, val_iou=0.446,
Train_Dice_CE_loss_epoch=0.868]

Metric Train_Dice_CE_loss improved by 0.004 >= min_delta = 0.0. New best score:
0.868

Epoch 14: 100%| | 144/144 [01:38<00:00, 1.46it/s, v_num=0,
Train_Dice_CE_loss_step=0.909, val_loss=0.871, val_dice=0.491, val_iou=0.446,
Train_Dice_CE_loss_epoch=0.868]TE 20

Epoch 14: 100%| | 144/144 [01:58<00:00, 1.22it/s, v_num=0,
Train_Dice_CE_loss_step=0.909, val_loss=0.867, val_dice=0.498, val_iou=0.453,
Train_Dice_CE_loss_epoch=0.863]

Metric Train_Dice_CE_loss improved by 0.005 >= min_delta = 0.0. New best score:
0.863

Epoch 15: 100%| | 144/144 [01:42<00:00, 1.40it/s, v_num=0,
Train_Dice_CE_loss_step=0.891, val_loss=0.867, val_dice=0.498, val_iou=0.453,
Train_Dice_CE_loss_epoch=0.863]TE 20

Epoch 15: 100%| | 144/144 [02:03<00:00, 1.17it/s, v_num=0,
Train_Dice_CE_loss_step=0.891, val_loss=0.859, val_dice=0.504, val_iou=0.455,
Train_Dice_CE_loss_epoch=0.850]

Metric Train_Dice_CE_loss improved by 0.013 >= min_delta = 0.0. New best score:
0.850

Epoch 16: 100%| | 144/144 [01:38<00:00, 1.46it/s, v_num=0,
Train_Dice_CE_loss_step=0.875, val_loss=0.859, val_dice=0.504, val_iou=0.455,
Train_Dice_CE_loss_epoch=0.850]TE 20

Epoch 16: 100%| | 144/144 [01:57<00:00, 1.22it/s, v_num=0,
Train_Dice_CE_loss_step=0.875, val_loss=0.856, val_dice=0.509, val_iou=0.460,
Train_Dice_CE_loss_epoch=0.846]

Metric Train_Dice_CE_loss improved by 0.004 >= min_delta = 0.0. New best score:
0.846

Epoch 17: 100%| | 144/144 [01:38<00:00, 1.47it/s, v_num=0,
Train_Dice_CE_loss_step=0.893, val_loss=0.856, val_dice=0.509, val_iou=0.460,
Train_Dice_CE_loss_epoch=0.846]TE 20

Epoch 17: 100%| | 144/144 [01:57<00:00, 1.22it/s, v_num=0,
Train_Dice_CE_loss_step=0.893, val_loss=0.855, val_dice=0.511, val_iou=0.462,
Train_Dice_CE_loss_epoch=0.842]

Metric Train_Dice_CE_loss improved by 0.004 >= min_delta = 0.0. New best score: 0.842

Epoch 18: 100%| | 144/144 [01:38<00:00, 1.47it/s, v_num=0,
Train_Dice_CE_loss_step=0.865, val_loss=0.855, val_dice=0.511, val_iou=0.462,
Train_Dice_CE_loss_epoch=0.842]TE 20
Epoch 18: 100%| | 144/144 [01:58<00:00, 1.22it/s, v_num=0,
Train_Dice_CE_loss_step=0.865, val_loss=0.850, val_dice=0.513, val_iou=0.463,
Train_Dice_CE_loss_epoch=0.838]

Metric Train_Dice_CE_loss improved by 0.003 >= min_delta = 0.0. New best score: 0.838

Epoch 19: 100%| | 144/144 [01:38<00:00, 1.46it/s, v_num=0,
Train_Dice_CE_loss_step=0.874, val_loss=0.850, val_dice=0.513, val_iou=0.463,
Train_Dice_CE_loss_epoch=0.838]TE 20

Logged image from epoch 19


Epoch 19: 100%| | 144/144 [01:57<00:00, 1.22it/s, v_num=0,
Train_Dice_CE_loss_step=0.874, val_loss=0.848, val_dice=0.514, val_iou=0.464,
Train_Dice_CE_loss_epoch=0.835]


Metric Train_Dice_CE_loss improved by 0.003 >= min_delta = 0.0. New best score: 0.835

`Trainer.fit` stopped: `max_epochs=20` reached.


Epoch 19: 100%| | 144/144 [01:58<00:00, 1.22it/s, v_num=0,
Train_Dice_CE_loss_step=0.874, val_loss=0.848, val_dice=0.514, val_iou=0.464,
Train_Dice_CE_loss_epoch=0.835]


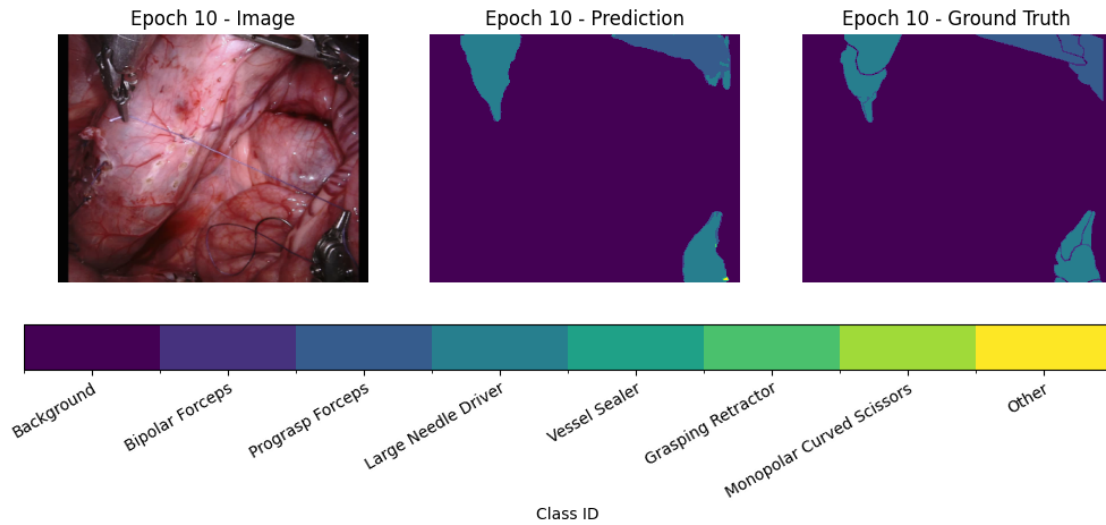
Training time: 39.52 minutes





Only 3 epochs recorded, plotting all.





```
[50]: if __name__ == "__main__":

    logger = TensorBoardLogger("tb_logs", name="part_instrument_seg")

    early_stop_callback = EarlyStopping(
        monitor="Train_Dice_CE_loss",           # metric name from self.log
        mode="min",                             # because lower loss is better
        patience=5,                             # epochs to wait before stopping
        verbose=True
    )
```

```

checkpoint_callback = ModelCheckpoint(
    monitor="Train_Dice_CE_loss",
    mode="min",
    save_top_k=1,
    dirpath="checkpoints/",
    filename="best-part-instrument-seg-basic-unet",
)

trainer = Trainer(
    accelerator="gpu",
    max_epochs=20,
    #limit_train_batches=0.1, # or 0.1 to use 10%
    logger=logger,
    callbacks=[early_stop_callback, checkpoint_callback],
)

start_train = time.time()
trainer.fit(
    model=part_instr_seg_basic_UNet_model,
    datamodule=part_instr_seg_endo_data
)
end_train = time.time()
print(f"Training time: {(end_train - start_train)/60:.2f} minutes")

# Plot the overlaid training and val loss curves per epoch
part_instr_seg_basic_UNet_model.plot_losses()

# Plot the IOU and DSC curves per epoch
part_instr_seg_basic_UNet_model.plot_metrics()

# Plot images from last epoch
part_instr_seg_basic_UNet_model.plot_result_by_epoch()

```

GPU available: True (cuda), used: True
 TPU available: False, using: 0 TPU cores
 HPU available: False, using: 0 HPUs
 Checkpoint directory C:\Users\dsumm\OneDrive\Documents\UMD ENPM Robotics
 Files\BIOE658B (Intro to Medical Image Analysis)\Project\code\checkpoints exists
 and is not empty.
 LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]

	Name	Type	Params	Mode	In sizes	Out sizes
0	model	BasicUNet	7.8 M	train	[1, 3, 256, 320]	[1, 21, 256, 320]
1	DICE_CE_Loss	DiceCELoss	0	train	?	?

```

-----
7.8 M      Trainable params
0          Non-trainable params
7.8 M      Total params
31.136     Total estimated model params size (MB)
143        Modules in train mode
0          Modules in eval mode

Sanity Checking: |          | 0/? [00:00<?, ?it/s]

The 'val_dataloader' does not have many workers which may be a bottleneck.
Consider increasing the value of the `num_workers` argument` to `num_workers=31`
in the `DataLoader` to improve performance.

The 'train_dataloader' does not have many workers which may be a bottleneck.
Consider increasing the value of the `num_workers` argument` to `num_workers=31`
in the `DataLoader` to improve performance.

Epoch 0: 100%|          | 144/144 [01:51<00:00, 1.29it/s, v_num=0,
Train_Dice_CE_loss_step=1.630]TE 20
Logged image from epoch 0
Epoch 0: 100%|          | 144/144 [02:15<00:00, 1.06it/s, v_num=0,
Train_Dice_CE_loss_step=1.630, val_loss=1.610, val_dice=0.0997, val_iou=0.0903,
Train_Dice_CE_loss_epoch=2.410]

Metric Train_Dice_CE_loss improved. New best score: 2.411

Epoch 1: 100%|          | 144/144 [01:41<00:00, 1.41it/s, v_num=0,
Train_Dice_CE_loss_step=1.420, val_loss=1.610, val_dice=0.0997, val_iou=0.0903,
Train_Dice_CE_loss_epoch=2.410]TE 20
Epoch 1: 100%|          | 144/144 [02:04<00:00, 1.16it/s, v_num=0,
Train_Dice_CE_loss_step=1.420, val_loss=1.400, val_dice=0.115, val_iou=0.107,
Train_Dice_CE_loss_epoch=1.470]

Metric Train_Dice_CE_loss improved by 0.943 >= min_delta = 0.0. New best score:
1.469

Epoch 2: 100%|          | 144/144 [01:41<00:00, 1.41it/s, v_num=0,
Train_Dice_CE_loss_step=1.380, val_loss=1.400, val_dice=0.115, val_iou=0.107,
Train_Dice_CE_loss_epoch=1.470]TE 20
Epoch 2: 100%|          | 144/144 [02:04<00:00, 1.15it/s, v_num=0,
Train_Dice_CE_loss_step=1.380, val_loss=1.330, val_dice=0.133, val_iou=0.120,
Train_Dice_CE_loss_epoch=1.360]

Metric Train_Dice_CE_loss improved by 0.106 >= min_delta = 0.0. New best score:
1.363

Epoch 3: 100%|          | 144/144 [01:42<00:00, 1.40it/s, v_num=0,
Train_Dice_CE_loss_step=1.290, val_loss=1.330, val_dice=0.133, val_iou=0.120,
Train_Dice_CE_loss_epoch=1.360]TE 20
Epoch 3: 100%|          | 144/144 [02:04<00:00, 1.16it/s, v_num=0,

```

Train_Dice_CE_loss_step=1.290, val_loss=1.290, val_dice=0.157, val_iou=0.138,
Train_Dice_CE_loss_epoch=1.310]

Metric Train_Dice_CE_loss improved by 0.053 >= min_delta = 0.0. New best score:
1.310

Epoch 4: 100%| | 144/144 [01:41<00:00, 1.42it/s, v_num=0,
Train_Dice_CE_loss_step=1.260, val_loss=1.290, val_dice=0.157, val_iou=0.138,
Train_Dice_CE_loss_epoch=1.310]TE 20

Epoch 4: 100%| | 144/144 [02:03<00:00, 1.17it/s, v_num=0,
Train_Dice_CE_loss_step=1.260, val_loss=1.260, val_dice=0.174, val_iou=0.152,
Train_Dice_CE_loss_epoch=1.270]

Metric Train_Dice_CE_loss improved by 0.037 >= min_delta = 0.0. New best score:
1.273

Epoch 5: 100%| | 144/144 [01:55<00:00, 1.25it/s, v_num=0,
Train_Dice_CE_loss_step=1.220, val_loss=1.260, val_dice=0.174, val_iou=0.152,
Train_Dice_CE_loss_epoch=1.270]TE 20

Epoch 5: 100%| | 144/144 [02:20<00:00, 1.02it/s, v_num=0,
Train_Dice_CE_loss_step=1.220, val_loss=1.190, val_dice=0.196, val_iou=0.169,
Train_Dice_CE_loss_epoch=1.220]

Metric Train_Dice_CE_loss improved by 0.058 >= min_delta = 0.0. New best score:
1.215

Epoch 6: 100%| | 144/144 [01:45<00:00, 1.37it/s, v_num=0,
Train_Dice_CE_loss_step=1.280, val_loss=1.190, val_dice=0.196, val_iou=0.169,
Train_Dice_CE_loss_epoch=1.220]TE 20

Epoch 6: 100%| | 144/144 [02:07<00:00, 1.13it/s, v_num=0,
Train_Dice_CE_loss_step=1.280, val_loss=1.180, val_dice=0.221, val_iou=0.191,
Train_Dice_CE_loss_epoch=1.190]

Metric Train_Dice_CE_loss improved by 0.029 >= min_delta = 0.0. New best score:
1.186

Epoch 7: 100%| | 144/144 [01:40<00:00, 1.43it/s, v_num=0,
Train_Dice_CE_loss_step=1.160, val_loss=1.180, val_dice=0.221, val_iou=0.191,
Train_Dice_CE_loss_epoch=1.190]TE 20

Epoch 7: 100%| | 144/144 [02:02<00:00, 1.17it/s, v_num=0,
Train_Dice_CE_loss_step=1.160, val_loss=1.140, val_dice=0.228, val_iou=0.196,
Train_Dice_CE_loss_epoch=1.160]

Metric Train_Dice_CE_loss improved by 0.027 >= min_delta = 0.0. New best score:
1.159

Epoch 8: 100%| | 144/144 [01:40<00:00, 1.44it/s, v_num=0,
Train_Dice_CE_loss_step=1.160, val_loss=1.140, val_dice=0.228, val_iou=0.196,
Train_Dice_CE_loss_epoch=1.160]TE 20

Epoch 8: 100%| | 144/144 [02:01<00:00, 1.19it/s, v_num=0,
Train_Dice_CE_loss_step=1.160, val_loss=1.120, val_dice=0.244, val_iou=0.211,
Train_Dice_CE_loss_epoch=1.140]

Metric Train_Dice_CE_loss improved by 0.024 >= min_delta = 0.0. New best score:
1.136

Epoch 9: 100%| | 144/144 [01:40<00:00, 1.44it/s, v_num=0,
Train_Dice_CE_loss_step=1.100, val_loss=1.120, val_dice=0.244, val_iou=0.211,
Train_Dice_CE_loss_epoch=1.140]TE 20
Epoch 9: 100%| | 144/144 [02:02<00:00, 1.18it/s, v_num=0,
Train_Dice_CE_loss_step=1.100, val_loss=1.100, val_dice=0.252, val_iou=0.219,
Train_Dice_CE_loss_epoch=1.110]

Metric Train_Dice_CE_loss improved by 0.027 >= min_delta = 0.0. New best score:
1.108

Epoch 10: 100%| | 144/144 [01:41<00:00, 1.42it/s, v_num=0,
Train_Dice_CE_loss_step=1.090, val_loss=1.100, val_dice=0.252, val_iou=0.219,
Train_Dice_CE_loss_epoch=1.110]TE 20
Logged image from epoch 10
Epoch 10: 100%| | 144/144 [02:03<00:00, 1.17it/s, v_num=0,
Train_Dice_CE_loss_step=1.090, val_loss=1.080, val_dice=0.262, val_iou=0.228,
Train_Dice_CE_loss_epoch=1.080]

Metric Train_Dice_CE_loss improved by 0.031 >= min_delta = 0.0. New best score:
1.077

Epoch 11: 100%| | 144/144 [01:40<00:00, 1.43it/s, v_num=0,
Train_Dice_CE_loss_step=1.070, val_loss=1.080, val_dice=0.262, val_iou=0.228,
Train_Dice_CE_loss_epoch=1.080]TE 20
Epoch 11: 100%| | 144/144 [02:01<00:00, 1.18it/s, v_num=0,
Train_Dice_CE_loss_step=1.070, val_loss=1.060, val_dice=0.268, val_iou=0.235,
Train_Dice_CE_loss_epoch=1.070]

Metric Train_Dice_CE_loss improved by 0.010 >= min_delta = 0.0. New best score:
1.067

Epoch 12: 100%| | 144/144 [01:40<00:00, 1.43it/s, v_num=0,
Train_Dice_CE_loss_step=1.060, val_loss=1.060, val_dice=0.268, val_iou=0.235,
Train_Dice_CE_loss_epoch=1.070]TE 20
Epoch 12: 100%| | 144/144 [02:01<00:00, 1.19it/s, v_num=0,
Train_Dice_CE_loss_step=1.060, val_loss=1.050, val_dice=0.273, val_iou=0.241,
Train_Dice_CE_loss_epoch=1.050]

Metric Train_Dice_CE_loss improved by 0.013 >= min_delta = 0.0. New best score:
1.054

Epoch 13: 100%| | 144/144 [01:32<00:00, 1.55it/s, v_num=0,
Train_Dice_CE_loss_step=1.070, val_loss=1.050, val_dice=0.273, val_iou=0.241,
Train_Dice_CE_loss_epoch=1.050]TE 20
Epoch 13: 100%| | 144/144 [01:53<00:00, 1.27it/s, v_num=0,
Train_Dice_CE_loss_step=1.070, val_loss=1.050, val_dice=0.277, val_iou=0.244,
Train_Dice_CE_loss_epoch=1.050]

Metric Train_Dice_CE_loss improved by 0.009 >= min_delta = 0.0. New best score:
1.045

Epoch 14: 100%| | 144/144 [01:32<00:00, 1.56it/s, v_num=0,
Train_Dice_CE_loss_step=1.060, val_loss=1.050, val_dice=0.277, val_iou=0.244,
Train_Dice_CE_loss_epoch=1.050]TE 20
Epoch 14: 100%| | 144/144 [01:52<00:00, 1.28it/s, v_num=0,
Train_Dice_CE_loss_step=1.060, val_loss=1.040, val_dice=0.282, val_iou=0.248,
Train_Dice_CE_loss_epoch=1.040]

Metric Train_Dice_CE_loss improved by 0.009 >= min_delta = 0.0. New best score:
1.037

Epoch 15: 100%| | 144/144 [01:32<00:00, 1.55it/s, v_num=0,
Train_Dice_CE_loss_step=1.070, val_loss=1.040, val_dice=0.282, val_iou=0.248,
Train_Dice_CE_loss_epoch=1.040]TE 20
Epoch 15: 100%| | 144/144 [01:53<00:00, 1.27it/s, v_num=0,
Train_Dice_CE_loss_step=1.070, val_loss=1.030, val_dice=0.285, val_iou=0.252,
Train_Dice_CE_loss_epoch=1.020]

Metric Train_Dice_CE_loss improved by 0.013 >= min_delta = 0.0. New best score:
1.024

Epoch 16: 100%| | 144/144 [01:32<00:00, 1.55it/s, v_num=0,
Train_Dice_CE_loss_step=1.030, val_loss=1.030, val_dice=0.285, val_iou=0.252,
Train_Dice_CE_loss_epoch=1.020]TE 20
Epoch 16: 100%| | 144/144 [01:53<00:00, 1.27it/s, v_num=0,
Train_Dice_CE_loss_step=1.030, val_loss=1.030, val_dice=0.290, val_iou=0.257,
Train_Dice_CE_loss_epoch=1.020]

Metric Train_Dice_CE_loss improved by 0.005 >= min_delta = 0.0. New best score:
1.019

Epoch 17: 100%| | 144/144 [01:32<00:00, 1.56it/s, v_num=0,
Train_Dice_CE_loss_step=1.030, val_loss=1.030, val_dice=0.290, val_iou=0.257,
Train_Dice_CE_loss_epoch=1.020]TE 20
Epoch 17: 100%| | 144/144 [01:52<00:00, 1.28it/s, v_num=0,
Train_Dice_CE_loss_step=1.030, val_loss=1.010, val_dice=0.290, val_iou=0.256,
Train_Dice_CE_loss_epoch=1.010]

Metric Train_Dice_CE_loss improved by 0.007 >= min_delta = 0.0. New best score:
1.012

Epoch 18: 100%| | 144/144 [01:32<00:00, 1.55it/s, v_num=0,
Train_Dice_CE_loss_step=1.030, val_loss=1.010, val_dice=0.290, val_iou=0.256,
Train_Dice_CE_loss_epoch=1.010]TE 20
Epoch 18: 100%| | 144/144 [01:53<00:00, 1.27it/s, v_num=0,
Train_Dice_CE_loss_step=1.030, val_loss=1.020, val_dice=0.295, val_iou=0.261,
Train_Dice_CE_loss_epoch=1.010]

Metric Train_Dice_CE_loss improved by 0.003 >= min_delta = 0.0. New best score:
1.009

Epoch 19: 100%| | 144/144 [01:32<00:00, 1.55it/s, v_num=0,
Train_Dice_CE_loss_step=1.030, val_loss=1.020, val_dice=0.295, val_iou=0.261,
Train_Dice_CE_loss_epoch=1.010]TE 20

Logged image from epoch 19


Epoch 19: 100%| | 144/144 [01:53<00:00, 1.27it/s, v_num=0,
Train_Dice_CE_loss_step=1.030, val_loss=1.020, val_dice=0.296, val_iou=0.263,
Train_Dice_CE_loss_epoch=1.010]


Metric Train_Dice_CE_loss improved by 0.002 >= min_delta = 0.0. New best score:
1.007

`Trainer.fit` stopped: `max_epochs=20` reached.

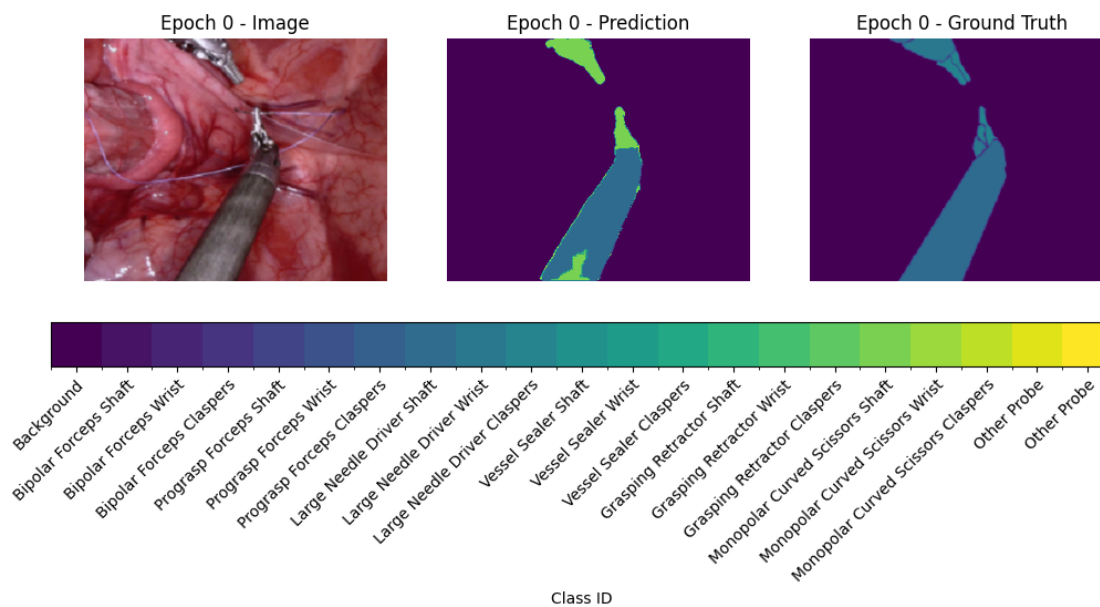
Epoch 19: 100%| | 144/144 [01:53<00:00, 1.27it/s, v_num=0,
Train_Dice_CE_loss_step=1.030, val_loss=1.020, val_dice=0.296, val_iou=0.263,
Train_Dice_CE_loss_epoch=1.010]

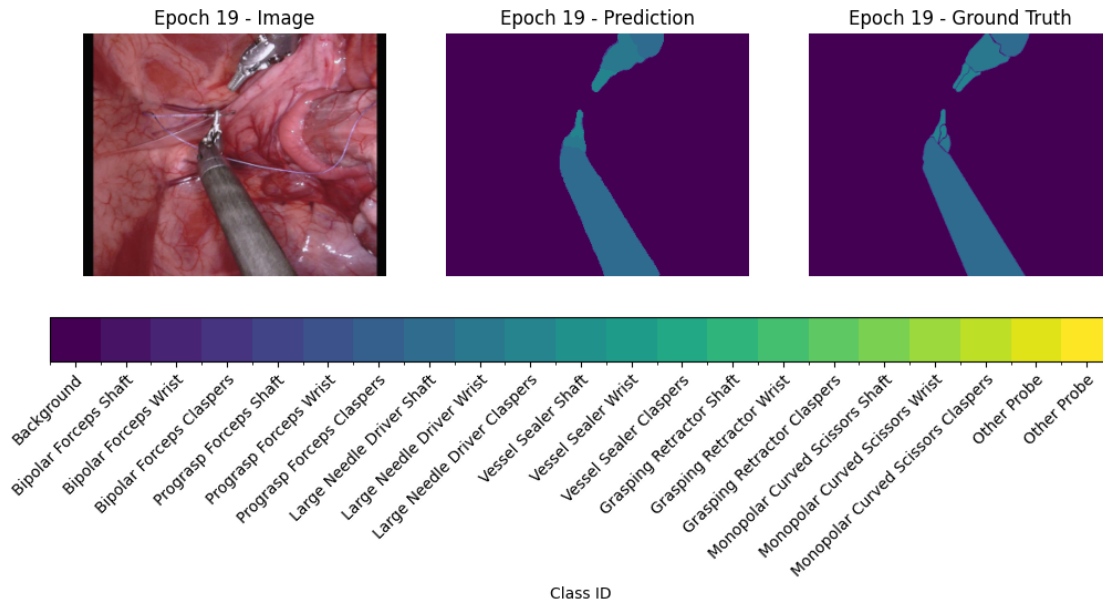
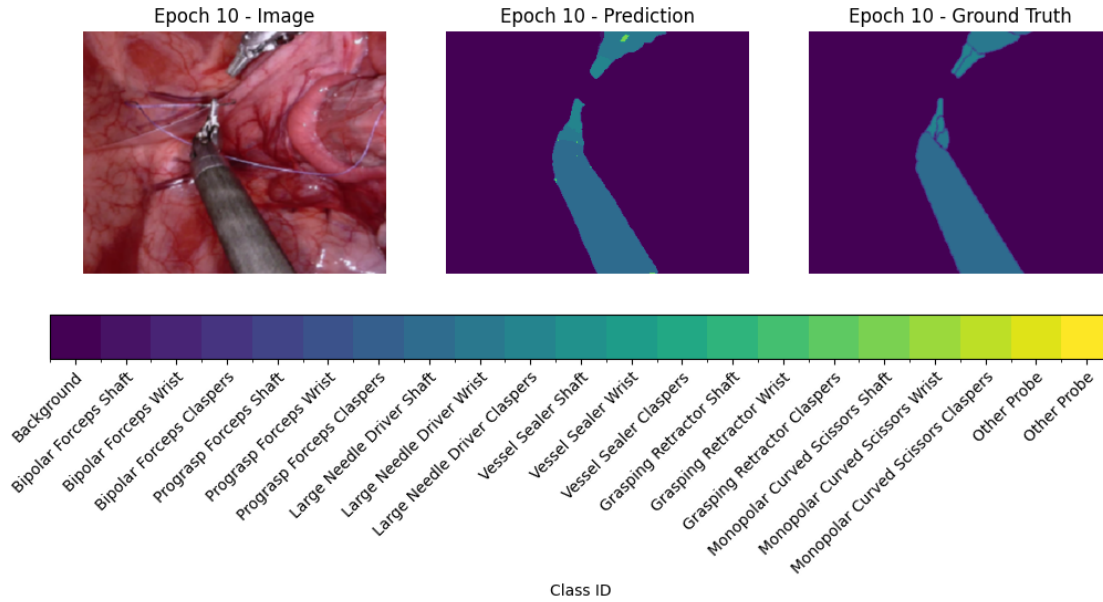
Training time: 40.55 minutes





Only 3 epochs recorded, plotting all.





```
[52]: os.makedirs('basic_UNet_final_models', exist_ok=True)

# Define file names with paths
binary_basic_UNet_model_filename = 'basicUNetmodels/binary_basic_UNet_model.pth'
part_seg_basic_UNet_model_filename = 'basicUNetmodels/part_seg_basic_UNet_model.
    ↪pth'
```

```

instr_seg_basic_UNet_model_filename = 'basicUNetmodels/
↳instr_seg_basic_UNet_model.pth'
part_instr_seg_basic_UNet_model_filename = 'basicUNetmodels/
↳part_instr_seg_basic_UNet_model.pth'

# Save the model parameters
torch.save(binary_basic_UNet_model.state_dict(),
↳binary_basic_UNet_model_filename)
torch.save(part_seg_basic_UNet_model.state_dict(),
↳part_seg_basic_UNet_model_filename)
torch.save(instr_seg_basic_UNet_model.state_dict(),
↳instr_seg_basic_UNet_model_filename)
torch.save(part_instr_seg_basic_UNet_model.state_dict(),
↳part_instr_seg_basic_UNet_model_filename)

print("Models saved in the 'basic_UNet_final_models' directory!")

```

Models saved in the 'basic_UNet_final_models' directory!