

## endo\_Basic\_Unet\_test\_mk1

May 13, 2025

```
[1]: import os
import numpy as np
from numpy.lib.stride_tricks import as_strided
import time
import matplotlib.pyplot as plt
from scipy.spatial.distance import directed_hausdorff

import torch
from torch.utils.data import DataLoader
from torch.utils.data import Dataset
from torch.utils.data import DataLoader, random_split
from torch.optim.lr_scheduler import StepLR

from pytorch_lightning import LightningDataModule
from pytorch_lightning import LightningModule
from pytorch_lightning import Trainer
from pytorch_lightning.callbacks import LearningRateMonitor, ModelCheckpoint
from pytorch_lightning.callbacks import EarlyStopping
from pytorch_lightning.loggers import TensorBoardLogger

from sklearn.model_selection import train_test_split
from sklearn.metrics import precision_score, recall_score, f1_score

from monai.networks.nets import BasicUNet
from monai.losses import DiceCELoss
from monai.metrics import DiceMetric, MeanIoU, HausdorffDistanceMetric,
    ConfusionMatrixMetric
from monai.transforms import (
    AsDiscreted,
    Compose,
    Resized,
    EnsureChannelFirstd,
    LoadImaged,
    ScaleIntensityd,
    ToTensord,
    RandFlipd,
    RandZoomd,
```

```

        ToTensord,
        AsDiscreted,
        CenterSpatialCropd
    )

```

C:\Users\dsumm\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.11\_qbz5n2kfra8p0\LocalCache\local-packages\Python311\site-packages\ignite\handlers\checkpoint.py:17: DeprecationWarning: `TorchScript` support for functional optimizers is deprecated and will be removed in a future PyTorch release. Consider using the `torch.compile` optimizer instead.

```

    from torch.distributed.optim import ZeroRedundancyOptimizer

```

```

[2]: # Custom dataset class for pytorch compatibility
# https://pytorch.org/tutorials/beginner/data_loading_tutorial.html
class EndoVis2017Dataset(Dataset):
    def __init__(self, label_subdir=None, test=False):
        self.data = []

        if label_subdir is None:
            raise ValueError("You must specify a `label_subdir` for ground_
↳ truth masks (e.g., 'instrument_seg_composite').")

        self.root_dir = "C:/Users/dsumm/OneDrive/Documents/UMD ENPM Robotics_
↳ Files/BIOE658B (Intro to Medical Image Analysis)/Project/dataset/test/"
        self.label_subdir = label_subdir

        # Recursively walk through directory to find left frame image paths and_
↳ GT image paths
        for subdir, dirs, files in os.walk(self.root_dir):
            if 'left_frames' in subdir:
                #print("Hit!")
                for file in sorted(files):
                    if file.endswith(('.png', '.jpg', '.jpeg')):
                        img_path = os.path.join(subdir, file)
                        #print(img_path)

                        gt_root = subdir.replace('left_frames', 'ground_truth')
                        mask_path = os.path.join(gt_root, self.label_subdir,
↳ file)

                        if os.path.exists(mask_path):
                            #print("Hit!")
                            self.data.append({"image": img_path, "label":
↳ mask_path}) # Dictionary for MONAI compatability

        if not test:
            transforms_list = [

```

```

        LoadImaged(keys=["image", "label"]), #
↳Loads image data and metadata from file path dictionaries
        EnsureChannelFirstd(keys=["image", "label"]), #
↳Adjust or add the channel dimension of input data to ensure channel_first
↳shape

        # Images are of nominal size 1280x1024 --> resizing for memory
↳efficiency
        CenterSpatialCropd(keys=["image", "label"], roi_size=(1024,
↳1280)), # Cropping background padding from images
        Resized(keys=["image", "label"], spatial_size=(256, 320)), #
↳Imported images are of various sizes: standardize to 320,256

        # Apply data augmentation techniques
        RandFlipd(keys=["image", "label"], prob=0.3, spatial_axis=1), #
↳# Horizontal axis flip imposed w/ 30% prob
        #RandRotate90d(keys=["image", "label"], prob=0.3, max_k=3), #
↳# Random 90° rotation imposed w/ 30% prob
        RandZoomd(keys=["image", "label"], prob=0.3, min_zoom=0.75,
↳max_zoom=1.25), # Zoom range (+/-25%) imposed w/ 30% prob
        #RandAdjustContrastd(keys=["image"], prob=0.3, gamma=(0.75, 1.
↳25)), # Contrast variation (+/-25%) imposed w/ 30% prob

        ScaleIntensityd(keys=["image"]), #
↳Scale the intensity of input image to the value range 0-1
        ToTensord(keys=["image", "label"]), #
↳Ensure data is of tensor type for pytorch usage
    ]
    else:
        transforms_list = [
            LoadImaged(keys=["image", "label"]), #
↳Loads image data and metadata from file path dictionaries
            EnsureChannelFirstd(keys=["image", "label"]), #
↳Adjust or add the channel dimension of input data to ensure channel_first
↳shape

            # Images are of nominal size 1280x1024 --> resizing for memory
↳efficiency
            CenterSpatialCropd(keys=["image", "label"], roi_size=(1024,
↳1280)), # Cropping background padding from images
            Resized(keys=["image", "label"], spatial_size=(256, 320)), #
↳Imported images are of various sizes: standardize to 320,256

            ScaleIntensityd(keys=["image"]), #
↳Scale the intensity of input image to the value range 0-1

```

```

        ToTensord(keys=["image", "label"]), #
↪ Ensure data is of tensor type for pytorch usage
    ]
    # Additional conditional transforms based on label_subdir
    if label_subdir == "binary_composite":
        transforms_list.append(AsDiscreted(keys=["label"], threshold=0.5))
↪    # Binary threshold for binary seg
    elif label_subdir == "part_seg_composite":
        transforms_list.append(AsDiscreted(keys=["label"], to_onehot=5))
↪    # 5 individual class labels for instrument independent part seg
    elif label_subdir == "TypeSegmentation":
        transforms_list.append(AsDiscreted(keys=["label"], to_onehot=8))
↪    # 8 individual class labels for part independent instrument seg
    elif label_subdir == "instrument_part_seg_composite":
        transforms_list.append(AsDiscreted(keys=["label"], to_onehot=21))
↪    # 26 individual class labels for instrument & part seg

    # Imposing MONAI transforms
    # https://docs.monai.io/en/stable/transforms.html
    self.transform = Compose(transforms_list)

    def __len__(self):
        # Returns number of imported samples
        length = len(self.data)
        return length

    def __getitem__(self, idx):
        # Return transformed sample from the dataset as dictated by the index
        sample = self.data[idx]
        return self.transform(sample)

```

```

[3]: class MONAIDataLoader(LightningDataModule):
    def __init__(self, dataset=None, batch_size: int = None, img_size: int =
↪ None, dimensions:int = None):
        super().__init__()
        if dataset is None:
            raise ValueError("No dataset given!")
        self.dataset = dataset
        self.test_dataset = dataset
        self.batch_size = batch_size
        self.pin_memory = True

        self.train, self.val = random_split(self.dataset, [
            int(len(self.dataset) * 0.8),
            len(self.dataset) - int(len(self.dataset) * 0.8)
        ])

```

```

        print(f"Train dataset size: {len(self.train)}")
        print(f"Validation dataset size: {len(self.val)}")
        print(f"Test dataset size: {len(self.test_dataset)}")

    def setup(self, stage=None):
        # required by PyTorch Lightning
        pass

    def train_dataloader(self):
        return DataLoader(self.train, batch_size=self.batch_size,
        ↪pin_memory=self.pin_memory)

    def val_dataloader(self):
        return DataLoader(self.val, batch_size=self.batch_size, pin_memory=self.
        ↪pin_memory)

    def test_dataloader(self):
        return DataLoader(self.test_dataset, batch_size=self.batch_size,
        ↪pin_memory=self.pin_memory)

```

```

[4]: class basic_UNet_Train(LightningModule):
    def __init__(self, img_size=(1, 3, 256, 320), batch_size=1, lr=0.001,
    ↪num_classes=1):
        super().__init__()

        self.save_hyperparameters()
        self.num_classes = num_classes
        print("num_classes", self.num_classes, num_classes, self.hparams.
        ↪num_classes)
        self.example_input_array = [torch.zeros(self.hparams.img_size)]

        self.test_step_outputs = [] # Initialize an empty list to store outputs

        self.dice_metric = DiceMetric(include_background=True,
        ↪reduction="mean", ignore_empty=True)
        self.iou_metric = MeanIoU(include_background=True, reduction="mean",
        ↪ignore_empty=True)
        self.hausdorff_metric = HausdorffDistanceMetric(
            include_background=True,
            distance_metric="euclidean",
            percentile=95,
            directed=False,
            reduction="mean"
        )
        self.confusion_metric = ConfusionMatrixMetric(
            metric_name=["precision", "recall", "f1 score"],
            include_background=False,

```

```

        compute_sample=False,
        reduction="mean"
    )

    # Metric tracking
    self.dice_scores = []
    self.iou_scores = []

    # Defining MONAI Unet model paramters
    self.model = BasicUNet(spatial_dims=2,          # 2D image so spatial dims
↪ = 2
                           in_channels=3,          # RGB input ultrasound image
                           out_channels=num_classes, # Binary
↪ segmentation mask output image
                           features= (32, 64, 128, 256, 512, 32), #
↪ standard Unet feature sizes (32, 32, 64, 128, 256, 32)
                           dropout=0.1)           # Dropout prob 10%

    # Using combined DICE and CE loss as loss function
    # Conditional loss function based on the number of classes
    if num_classes == 1:
        self.DICE_CE_Loss = DiceCELoss(
            include_background=False, # Exclude background class
            sigmoid=True, # Use softmax for multiclass segmentation
            softmax=False, # Apply softmax for multiclass
            lambda_dice=1.0, # Adjust the weight for Dice loss
            lambda_ce=1.0, # Adjust the weight for Cross-Entropy loss
            reduction='mean' # Use mean reduction
        )
    else:
        self.DICE_CE_Loss = DiceCELoss(
            include_background=False, # Exclude background class
            sigmoid=False, # Use softmax for multiclass segmentation
            softmax=True, # Apply softmax for multiclass
            lambda_dice=1.0, # Adjust the weight for Dice loss
            lambda_ce=1.0, # Adjust the weight for Cross-Entropy loss
            reduction='mean' # Use mean reduction
        )

    # For storing images for the last epoch
    self.last_image = []
    self.last_pred = []
    self.last_mask = []
    self.logged_epochs = []

    def forward(self, inputs):
        outputs = self.model(inputs)

```

```

return outputs

def test_step(self, batch, batch_idx):
    # Prepare input and ground truth
    inputs, gt_input = self._prepare_batch(batch)
    outputs = self.forward(inputs)

    if self.hparams.num_classes == 1:
        # Binary segmentation
        probs = torch.sigmoid(outputs)
        preds = (probs > 0.5).float()
        gt_input = (gt_input > 0.5).float()

    else:
        probs = torch.softmax(outputs, dim=1)
        preds = torch.nn.functional.one_hot(torch.argmax(probs, dim=1),
↪num_classes=self.num_classes)
        preds = preds.permute(0, 3, 1, 2).float() # Shape: [B, C, H, W]

    # MONAI metrics
    self.dice_metric(y_pred=preds, y=gt_input)
    self.iou_metric(y_pred=preds, y=gt_input)

    # Hausdorff: safe only per image if non-empty
    for i in range(preds.shape[0]):
        pred_i = preds[i]
        gt_i = gt_input[i]
        if torch.any(pred_i) and torch.any(gt_i): # Check both non-empty
            self.hausdorff_metric(y_pred=pred_i.unsqueeze(0), y=gt_i.
↪unsqueeze(0))
        else:
            print(f"[Info] Skipping HD metric for empty prediction or GT in
↪batch index {i}")
            #self.hausdorff_metric(y_pred=preds, y=gt_input)
            self.confusion_metric(y_pred=preds, y=gt_input)

    # Extract Dice, IoU, Hausdorff from MONAI
    # Aggregate & safely handle NaNs
    dice = torch.nan_to_num(self.dice_metric.aggregate(), nan=0.0).item()
    iou = torch.nan_to_num(self.iou_metric.aggregate(), nan=0.0).item()
    hausdorff = torch.nan_to_num(self.hausdorff_metric.aggregate(), nan=0.
↪0).item()
    #hausdorff = self.hausdorff_metric.aggregate().item()
    #hausdorff = float('nan') if torch.isnan(torch.tensor(hausdorff)) else
↪hausdorff
    #hausdorff = torch.nan_to_num(hausdorff, nan=0.0)

```

```

self.dice_metric.reset()
self.iou_metric.reset()
self.hausdorff_metric.reset()

# Extract precision, recall, f1 score
confusion_metrics = self.confusion_metric.aggregate()
precision, recall, f1 = [m.item() for m in confusion_metrics]
self.confusion_metric.reset()

# Log metrics
self.log("test_dice", dice, prog_bar=True)
self.log("test_iou", iou, prog_bar=True)
self.log("test_hausdorff", hausdorff, prog_bar=True)
self.log("test_precision", precision, prog_bar=True)
self.log("test_recall", recall, prog_bar=True)
self.log("test_f1", f1, prog_bar=True)

# Return for aggregation
out = {
    "test_dice": torch.tensor(dice),
    "test_iou": torch.tensor(iou),
    "test_precision": torch.tensor(precision),
    "test_recall": torch.tensor(recall),
    "test_f1": torch.tensor(f1),
    "test_hausdorff": torch.tensor(hausdorff)
}

self.test_step_outputs.append(out)
return out

def on_test_epoch_end(self):
    # Aggregate the results across all batches in the epoch
    avg_dice = torch.stack([x["test_dice"] for x in self.
↪test_step_outputs]).mean()
    avg_iou = torch.stack([x["test_iou"] for x in self.test_step_outputs]).
↪mean()
    avg_hausdorff = torch.stack([x["test_hausdorff"] for x in self.
↪test_step_outputs]).mean()
    avg_precision = torch.stack([x["test_precision"] for x in self.
↪test_step_outputs]).mean()
    avg_recall = torch.stack([x["test_recall"] for x in self.
↪test_step_outputs]).mean()
    avg_f1 = torch.stack([x["test_f1"] for x in self.test_step_outputs]).
↪mean()

    print(f"\n Test Metrics:")

```



```

        f"\n    Dice      : {avg_dice.item():.4f}"
        f"\n    IoU       : {avg_iou.item():.4f}"
        f"\n    Hausdorff : {avg_hausdorff.item():.4f}"
        f"\n    Precision : {avg_precision.item():.4f}"
        f"\n    Recall    : {avg_recall.item():.4f}"
        f"\n    F1 Score  : {avg_f1.item():.4f}"

    # Clear for next epoch
    self.test_step_outputs.clear()

def _prepare_batch(self, batch):
    return batch['image'], batch['label']

def training_step(self, batch, batch_idx):
    inputs, gt_input = self._prepare_batch(batch)
    outputs = self.forward(inputs)
    loss = self.DICE_CE_Loss(outputs, gt_input)
    self.log(f"Train_Dice_CE_loss", loss, on_epoch=True, prog_bar=True)
    if batch_idx == len(batch) - 1:
        self.train_losses.append(loss.item())

    return loss

def validation_step(self, batch, batch_idx):

    # Gets labels for input and corresponding ground truth
    inputs, gt_input = self._prepare_batch(batch)
    outputs = self.forward(inputs)
    loss = self.DICE_CE_Loss(outputs, gt_input)
    self.log("val_loss", loss, on_step=False, on_epoch=True, prog_bar=True)

    if self.hparams.num_classes == 1:
        probs = torch.sigmoid(outputs)
        preds = (probs > 0.5).float()
        # Ensure ground truth is binary (i.e., 0 or 1)
        gt_input = (gt_input > 0.5).float() # Threshold the ground truth
    ↪if needed

    intersection = (preds * gt_input).sum()
    union = preds.sum() + gt_input.sum()
    bin_dice_score = 2.0 * intersection / (union + 1e-8) # Avoid
    ↪division by zero
    # IoU score calculation for binary segmentation
    bin_iou_score = intersection / (union - intersection + 1e-8) #
    ↪Avoid division by zero

```

```

        self.log("val_dice", bin_dice_score, on_step=False, on_epoch=True,
        ↪prog_bar=True)
        self.log("val_iou", bin_iou_score, on_step=False, on_epoch=True,
        ↪prog_bar=True)

    else:
        probs = torch.softmax(outputs, dim=1)
        preds = torch.nn.functional.one_hot(torch.argmax(probs, dim=1),
        ↪num_classes=self.num_classes)
        preds = preds.permute(0, 3, 1, 2).float() # Shape: [B, C, H, W]

        self.dice_metric(y_pred=preds, y=gt_input)
        self.iou_metric(y_pred=preds, y=gt_input)

    if self.trainer.sanity_checking:
        return # skip logging during sanity check

    # Append validation loss at the end of each epoch
    if batch_idx == len(batch) - 1:
        self.val_losses.append(loss.item())

    # For binary segmentation: apply sigmoid and threshold
    if self.hparams.num_classes == 1:
        outputs = torch.sigmoid(outputs)
        outputs = (outputs > 0.5).float() # Convert probabilities to
        ↪binary mask

        self.dice_scores.append(bin_dice_score)
        self.iou_scores.append(bin_iou_score)

    # For multiclass segmentation: apply softmax
    else:
        outputs = torch.softmax(outputs, dim=1) # Apply softmax for
        ↪multi-class outputs
        dice = self.dice_metric.aggregate()[0].item()
        #print("Dice", dice)
        iou = self.iou_metric.aggregate()[0].item()
        #print("IOU", iou)
        self.dice_metric.reset()
        self.iou_metric.reset()
        self.dice_scores.append(dice)
        self.iou_scores.append(iou)
        self.log("val_dice", dice, on_step=False, on_epoch=True,
        ↪prog_bar=True)
        self.log("val_iou", iou, on_step=False, on_epoch=True,
        ↪prog_bar=True)

```

```

        # Normalize and convert tensor to 3 channels (RGB) for visualization
    def process(last):
        # Detach from cpu to not interrupt training
        # https://stackoverflow.com/questions/63582590/
        ↪why-do-we-call-detach-before-calling-numpy-on-a-pytorch-tensor
        last = last[0].detach().cpu()

        # Min max normalization
        # https://www.codecademy.com/article/normalization
        last = (last - last.min()) / (last.max() - last.min() + 1e-8)

        # If grayscale, reshape last image to RGB for display by ↵
        ↪replicating gray value twice
        # https://discuss.pytorch.org/t/convert-grayscale-images-to-rgb/
        ↪113422
        return last.repeat(3, 1, 1) if last.shape[0] == 1 else last

    current_epoch = self.current_epoch
    total_epochs = self.trainer.max_epochs
    print("TE", total_epochs)

    if current_epoch == 0 or current_epoch == total_epochs - 1 or ↵
    ↪current_epoch == total_epochs // 2:
        self.last_image.append(process(inputs))
        self.last_pred.append(process(outputs))
        self.last_mask.append(process(gt_input))
        self.logged_epochs.append(current_epoch)
        print(f"Logged image from epoch {current_epoch}")

    return loss

    def plot_losses(self):
        min_len = min(len(self.train_losses), len(self.val_losses))
        epochs = range(1, min_len + 1)

        # Plotting training vs validation loss
        plt.figure(figsize=(10, 6))
        plt.plot(epochs, self.train_losses[:len(epochs)], label="Training ↵
        ↪Loss", color='blue')
        plt.plot(epochs, self.val_losses[:len(epochs)], label="Validation ↵
        ↪Loss", color='orange')
        plt.title("Training vs Validation Loss")
        plt.xlabel("Epochs")
        plt.ylabel("Loss")
        plt.legend()
        plt.show()

```

```

def plot_metrics(self):
    epochs = range(1, len(self.dice_scores) + 1)

    # Convert to CPU floats if necessary
    dice = [d.cpu().item() if torch.is_tensor(d) else d for d in self.
↪dice_scores]
    iou = [i.cpu().item() if torch.is_tensor(i) else i for i in self.
↪iou_scores]

    plt.figure(figsize=(10, 6))
    plt.plot(epochs, dice, label='Dice Coefficient')
    plt.plot(epochs, iou, label='IoU')
    plt.xlabel("Epochs")
    plt.ylabel("Score")
    plt.title("Validation Metrics Over Time")
    plt.legend()
    plt.show()

def plot_result_by_epoch(self):
    total_epochs = len(self.last_image)

    if total_epochs < 5:
        print(f"Only {total_epochs} epochs recorded, plotting all.")
        selected_epochs = list(range(total_epochs))
    else:
        print(f"{total_epochs} epochs recorded, bug in code.")

    for epoch_idx in selected_epochs:
        epoch_num = self.logged_epochs[epoch_idx] if hasattr(self,
↪"logged_epochs") else epoch_idx
        img = self.last_image[epoch_idx]
        pred = self.last_pred[epoch_idx]
        mask = self.last_mask[epoch_idx]

        fig, ax = plt.subplots(1, 3, figsize=(12, 4))

        ax[0].imshow(np.transpose(img.numpy(), (1, 2, 0)))
        ax[0].set_title(f"Epoch {epoch_num} - Image")
        ax[0].axis("off")

        if self.hparams.num_classes == 1:
            ax[1].imshow(np.transpose(pred.numpy(), (1, 2, 0)))
            ax[1].set_title(f"Epoch {epoch_num} - Prediction")
            ax[1].axis("off")

        ax[2].imshow(np.transpose(mask.numpy(), (1, 2, 0)))

```

```

ax[2].set_title(f"Epoch {epoch_num} - Ground Truth")
ax[2].axis("off")
else:
    # Define the colormap and normalization
    num_classes = self.hparams.num_classes
    cmap = plt.get_cmap('viridis', num_classes)
    bounds = np.arange(num_classes + 1) - 0.5
    norm = plt.matplotlib.colors.BoundaryNorm(bounds, cmap.N)

    # Convert one-hot encoded predictions and masks to
    ↪single-channel class labels
    pred_mask = torch.argmax(pred, dim=0).cpu().numpy()
    true_mask = torch.argmax(mask, dim=0).cpu().numpy()

    # Apply consistent colormap and normalization
    im1 = ax[1].imshow(pred_mask, cmap=cmap, norm=norm)
    ax[1].set_title(f"Epoch {epoch_num} - Prediction")
    ax[1].axis("off")

    im2 = ax[2].imshow(true_mask, cmap=cmap, norm=norm)
    ax[2].set_title(f"Epoch {epoch_num} - Ground Truth")
    ax[2].axis("off")

    im_for_cbar = im1 # just need one mappable

    # Adjust layout to leave space at the bottom
    fig.subplots_adjust(bottom=0.25) # tweak this if labels get cut
    ↪off

    # Add a new axis below the plots for the colorbar
    cbar_ax = fig.add_axes([0.1, 0.1, 0.8, 0.10]) # [left, bottom,
    ↪width, height]
    cbar = fig.colorbar(im_for_cbar, cax=cbar_ax,
    ↪orientation='horizontal', ticks=np.arange(num_classes))

    # Add colorbar below the plots
    #cbar = fig.colorbar(im1, ax=ax.ravel().tolist(),
    ↪orientation='horizontal',
    #ticks=np.arange(num_classes), pad=0.15, fraction=0.05)

    # Set class labels
    if num_classes == 5:
        cbar.ax.set_xticklabels(['Background', 'Shaft', 'Wrist',
        ↪'Claspers', 'Probe'])
    elif num_classes == 8:

```

```

        cbar.ax.set_xticklabels(['Background', 'Bipolar Forceps',
↪ 'Prograsp Forceps', 'Large Needle Driver',
                                'Vessel Sealer', 'Grasping
↪ Retractor', 'Monopolar Curved Scissors', 'Other'])

        plt.setp(cbar.ax.get_xticklabels(), rotation=30,
↪ ha="right", rotation_mode="anchor")
        elif num_classes == 21:
            cbar.ax.set_xticklabels([
                "Background",
                "Bipolar Forceps Shaft", "Bipolar Forceps Wrist",
↪ "Bipolar Forceps Claspers",
                "Prograsp Forceps Shaft", "Prograsp Forceps Wrist",
↪ "Prograsp Forceps Claspers",
                "Large Needle Driver Shaft", "Large Needle Driver
↪ Wrist", "Large Needle Driver Claspers",
                "Vessel Sealer Shaft", "Vessel Sealer Wrist", "Vessel
↪ Sealer Claspers",
                "Grasping Retractor Shaft", "Grasping Retractor Wrist",
↪ "Grasping Retractor Claspers",
                "Monopolar Curved Scissors Shaft", "Monopolar Curved
↪ Scissors Wrist", "Monopolar Curved Scissors Claspers",
                "Other Probe", "Other Probe"
            ])
        plt.setp(cbar.ax.get_xticklabels(), rotation=45,
↪ ha="right", rotation_mode="anchor")

        cbar.set_label('Class ID')

plt.show()

```

```

[5]: binary_basic_UNet_model = basic_UNet_Train(num_classes=1)
binary_basic_UNet_model.load_state_dict(torch.load('C:/Users/dsumm/OneDrive/
↪ Documents/UMD ENPM Robotics Files/BIOE658B (Intro to Medical Image Analysis)/
↪ Project/results/Basic_UNet/basicUNetmodels/binary_basic_UNet_model.pth'))

binary_endo_images = EndoVis2017Dataset(label_subdir='binarySegmentation',
↪ test=True)
binary_endo_data = MONAIDataLoader(dataset=binary_endo_images, batch_size=20)

trainer = Trainer(accelerator="gpu", devices=1)
trainer.test(model=binary_basic_UNet_model, datamodule=binary_endo_data)

```

num\_classes 1 1 1

BasicUNet features: (32, 64, 128, 256, 512, 32).

You are using `torch.load` with `weights\_only=False` (the current default

value), which uses the default pickle module implicitly. It is possible to construct malicious pickle data which will execute arbitrary code during unpickling (See <https://github.com/pytorch/pytorch/blob/main/SECURITY.md#untrusted-models> for more details). In a future release, the default value for `weights\_only` will be flipped to `True`. This limits the functions that could be executed during unpickling. Arbitrary objects will no longer be allowed to be loaded via this mode unless they are explicitly allowlisted by the user via `torch.serialization.add\_safe\_globals`. We recommend you start setting `weights\_only=True` for any use case where you don't have full control of the loaded file. Please open an issue on GitHub for any issues related to this experimental feature.

Using default `ModelCheckpoint`. Consider installing `litmodels` package to enable `LitModelCheckpoint` for automatic upload to the Lightning model registry.

GPU available: True (cuda), used: True

TPU available: False, using: 0 TPU cores

HPU available: False, using: 0 HPUs

You are using a CUDA device ('NVIDIA GeForce RTX 4070 Laptop GPU') that has Tensor Cores. To properly utilize them, you should set

`torch.set\_float32\_matmul\_precision('medium' | 'high')` which will trade-off precision for performance. For more details, read [https://pytorch.org/docs/stable/generated/torch.set\\_float32\\_matmul\\_precision.html#torch.set\\_float32\\_matmul\\_precision](https://pytorch.org/docs/stable/generated/torch.set_float32_matmul_precision.html#torch.set_float32_matmul_precision)

LOCAL\_RANK: 0 - CUDA\_VISIBLE\_DEVICES: [0]

Train dataset size: 720

Validation dataset size: 180

Test dataset size: 900

The 'test\_dataloader' does not have many workers which may be a bottleneck.

Consider increasing the value of the `num\_workers` argument` to `num\_workers=31` in the `DataLoader` to improve performance.

Testing: | | 0/? [00:00<?, ?it/s]

Test Metrics:

Dice : 0.9100  
IoU : 0.8552  
Hausdorff : 28.7077  
Precision : 0.8954  
Recall : 0.9532  
F1 Score : 0.9139

Test metric	DataLoader 0
test_dice	0.9099989533424377
test_f1	0.9138563275337219

```

test_hausdorff      28.707733154296875
test_iou            0.8551986813545227
test_precision      0.895444393157959
test_recall         0.9531691670417786

```

```

[5]: [{'test_dice': 0.9099989533424377,
      'test_iou': 0.8551986813545227,
      'test_hausdorff': 28.707733154296875,
      'test_precision': 0.895444393157959,
      'test_recall': 0.9531691670417786,
      'test_f1': 0.9138563275337219}]

```

```

[6]: binary_basic_UNet_model.eval().cuda()  # <<< This is important
N_BATCHES = 10  # Set number of batches to evaluate
times = []

with torch.no_grad():
    for i, batch in enumerate(binary_endo_data.test_dataloader()):
        if i >= N_BATCHES:
            break
        inputs = batch["image"].cuda()
        start_time = time.time()
        outputs = binary_basic_UNet_model(inputs)
        torch.cuda.synchronize()  # Ensures accurate timing on GPU
        end_time = time.time()
        times.append(end_time - start_time)

avg_infer_time = np.mean(times) / inputs.shape[0]  # Per image
print(f"Average inference time per image over {N_BATCHES * inputs.shape[0]}
      ↪images: {avg_infer_time:.6f} seconds")

```

Average inference time per image over 200 images: 0.007365 seconds

```

[7]: part_seg_basic_UNet_model = basic_UNet_Train(num_classes=5)
part_seg_basic_UNet_model.load_state_dict(torch.load('C:/Users/dsumm/OneDrive/
      ↪Documents/UMD ENPM Robotics Files/BIOE658B (Intro to Medical Image Analysis)/
      ↪Project/results/Basic_UNet/basicUNetmodels/part_seg_basic_UNet_model.pth'))

part_seg_endo_images = EndoVis2017Dataset(label_subdir='part_seg_composite',
      ↪test=True)
part_seg_endo_data = MONAIDataLoader(dataset=part_seg_endo_images,
      ↪batch_size=10)

trainer = Trainer(accelerator="gpu", devices=1)
trainer.test(model=part_seg_basic_UNet_model, datamodule=part_seg_endo_data)

```

You are using `torch.load` with `weights\_only=False` (the current default



value), which uses the default pickle module implicitly. It is possible to construct malicious pickle data which will execute arbitrary code during unpickling (See <https://github.com/pytorch/pytorch/blob/main/SECURITY.md#untrusted-models> for more details). In a future release, the default value for `weights\_only` will be flipped to `True`. This limits the functions that could be executed during unpickling. Arbitrary objects will no longer be allowed to be loaded via this mode unless they are explicitly allowlisted by the user via `torch.serialization.add\_safe\_globals`. We recommend you start setting `weights\_only=True` for any use case where you don't have full control of the loaded file. Please open an issue on GitHub for any issues related to this experimental feature.

Using default `ModelCheckpoint`. Consider installing `litmodels` package to enable `LitModelCheckpoint` for automatic upload to the Lightning model registry.

GPU available: True (cuda), used: True

TPU available: False, using: 0 TPU cores

HPU available: False, using: 0 HPUs

LOCAL\_RANK: 0 - CUDA\_VISIBLE\_DEVICES: [0]

num\_classes 5 5 5

BasicUNet features: (32, 64, 128, 256, 512, 32).

Train dataset size: 720

Validation dataset size: 180

Test dataset size: 900

Testing: | | 0/? [00:00<?, ?it/s]

the ground truth of class 4 is all 0, this may result in nan/inf distance.

the prediction of class 4 is all 0, this may result in nan/inf distance.

the ground truth of class 3 is all 0, this may result in nan/inf distance.

the prediction of class 2 is all 0, this may result in nan/inf distance.

the prediction of class 3 is all 0, this may result in nan/inf distance.

the ground truth of class 2 is all 0, this may result in nan/inf distance.

Test Metrics:

Dice : 0.7990

IoU : 0.7129

Hausdorff : 43.1655

Precision : 0.7664

Recall : 0.8284

F1 Score : 0.7872

Test metric	DataLoader 0
test_dice	0.799024224281311
test_f1	0.787174642086029
test_hausdorff	43.165470123291016

```

test_iou          0.7128857374191284
test_precision    0.7663677334785461
test_recall       0.8283786773681641

```

```

[7]: [{'test_dice': 0.799024224281311,
      'test_iou': 0.7128857374191284,
      'test_hausdorff': 43.165470123291016,
      'test_precision': 0.7663677334785461,
      'test_recall': 0.8283786773681641,
      'test_f1': 0.787174642086029}]

```

```

[8]: part_seg_basic_UNet_model.eval().cuda() # <<< This is important
N_BATCHES = 20 # Set number of batches to evaluate
times = []

with torch.no_grad():
    for i, batch in enumerate(part_seg_endo_data.test_dataloader()):
        if i >= N_BATCHES:
            break
        inputs = batch["image"].cuda()
        start_time = time.time()
        outputs = part_seg_basic_UNet_model(inputs)
        torch.cuda.synchronize() # Ensures accurate timing on GPU
        end_time = time.time()
        times.append(end_time - start_time)

avg_infer_time = np.mean(times) / inputs.shape[0] # Per image
print(f"Average inference time per image over {N_BATCHES * inputs.shape[0]}
      ↳images: {avg_infer_time:.6f} seconds")

```

Average inference time per image over 200 images: 0.007161 seconds

```

[9]: instr_seg_basic_UNet_model = basic_UNet_Train(num_classes=8)
instr_seg_basic_UNet_model.load_state_dict(torch.load('C:/Users/dsumm/OneDrive/
↳Documents/UMD ENPM Robotics Files/BI0E658B (Intro to Medical Image Analysis)/
↳Project/results/Basic_UNet/basicUNetmodels/instr_seg_basic_UNet_model.pth'))

instr_seg_endo_images = EndoVis2017Dataset(label_subdir='TypeSegmentation',
↳test=True)
instr_seg_endo_data = MONAIDataLoader(dataset=instr_seg_endo_images,
↳batch_size=5)

trainer = Trainer(accelerator="gpu", devices=1)
trainer.test(model=instr_seg_basic_UNet_model, datamodule=instr_seg_endo_data)

```

You are using `torch.load` with `weights\_only=False` (the current default value), which uses the default pickle module implicitly. It is possible to

construct malicious pickle data which will execute arbitrary code during unpickling (See <https://github.com/pytorch/pytorch/blob/main/SECURITY.md#untrusted-models> for more details). In a future release, the default value for `weights\_only` will be flipped to `True`. This limits the functions that could be executed during unpickling. Arbitrary objects will no longer be allowed to be loaded via this mode unless they are explicitly allowlisted by the user via `torch.serialization.add\_safe\_globals`. We recommend you start setting `weights\_only=True` for any use case where you don't have full control of the loaded file. Please open an issue on GitHub for any issues related to this experimental feature.

Using default `ModelCheckpoint`. Consider installing `litmodels` package to enable `LitModelCheckpoint` for automatic upload to the Lightning model registry.

```
num_classes 8 8 8
BasicUNet features: (32, 64, 128, 256, 512, 32).
Train dataset size: 720
Validation dataset size: 180
Test dataset size: 900
```

```
GPU available: True (cuda), used: True
TPU available: False, using: 0 TPU cores
HPU available: False, using: 0 HPUs
LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]
```

```
Testing: |          | 0/? [00:00<?, ?it/s]
```

```
the ground truth of class 5 is all 0, this may result in nan/inf distance.
the ground truth of class 6 is all 0, this may result in nan/inf distance.
the ground truth of class 7 is all 0, this may result in nan/inf distance.
the prediction of class 6 is all 0, this may result in nan/inf distance.
the prediction of class 5 is all 0, this may result in nan/inf distance.
the prediction of class 7 is all 0, this may result in nan/inf distance.
the prediction of class 1 is all 0, this may result in nan/inf distance.
```

Test Metrics:

```
Dice      : 0.5215
IoU       : 0.4717
Hausdorff : 47.7621
Precision : 0.6344
Recall    : 0.6564
F1 Score  : 0.6392
```

Test metric	DataLoader 0
test_dice	0.5215451717376709
test_f1	0.6391516327857971
test_hausdorff	47.76210021972656

```

test_iou          0.4717252850532532
test_precision    0.6343963146209717
test_recall       0.6564075350761414

```

```

[9]: [{'test_dice': 0.5215451717376709,
      'test_iou': 0.4717252850532532,
      'test_hausdorff': 47.76210021972656,
      'test_precision': 0.6343963146209717,
      'test_recall': 0.6564075350761414,
      'test_f1': 0.6391516327857971}]

```

```

[10]: instr_seg_basic_UNet_model.eval().cuda() # <<< This is important
N_BATCHES = 40 # Set number of batches to evaluate
times = []

with torch.no_grad():
    for i, batch in enumerate(instr_seg_endo_data.test_dataloader()):
        if i >= N_BATCHES:
            break
        inputs = batch["image"].cuda()
        start_time = time.time()
        outputs = instr_seg_basic_UNet_model(inputs)
        torch.cuda.synchronize() # Ensures accurate timing on GPU
        end_time = time.time()
        times.append(end_time - start_time)

avg_infer_time = np.mean(times) / inputs.shape[0] # Per image
print(f"Average inference time per image over {N_BATCHES * inputs.shape[0]}␣
↪images: {avg_infer_time:.6f} seconds")

```

Average inference time per image over 200 images: 0.248426 seconds