

endo_SwinUNETR_mk1

May 12, 2025

```
[18]: import os
import numpy as np
from numpy.lib.stride_tricks import as_strided
import time
import matplotlib.pyplot as plt
from scipy.spatial.distance import directed_hausdorff
from tqdm.notebook import tqdm

import torch
from torch.utils.data import DataLoader
from torch.utils.data import Dataset
from torch.utils.data import DataLoader, random_split
from torch.optim.lr_scheduler import StepLR

from pytorch_lightning import LightningDataModule
from pytorch_lightning import LightningModule
from pytorch_lightning import Trainer
from pytorch_lightning.callbacks import LearningRateMonitor, ModelCheckpoint
from pytorch_lightning.callbacks import EarlyStopping
from pytorch_lightning.loggers import TensorBoardLogger

from sklearn.model_selection import train_test_split

from monai.networks.nets import SwinUNETR
from monai.losses import DiceCELoss
from monai.metrics import DiceMetric, MeanIoU, HausdorffDistanceMetric
from monai.transforms import (
    AsDiscreted,
    Compose,
    Resized,
    EnsureChannelFirstd,
    LoadImaged,
    ScaleIntensityd,
    ToTensord,
    RandFlipd,
    RandZoomd,
    ToTensord,
```

```

        AsDiscreted,
        CenterSpatialCropd
    )

```

```

[19]: # Custom dataset class for pytorch compatibility
# https://pytorch.org/tutorials/beginner/data_loading_tutorial.html
class EndoVis2017Dataset(Dataset):
    def __init__(self, label_subdir=None):
        self.data = []

        if label_subdir is None:
            raise ValueError("You must specify a `label_subdir` for ground_
↳truth masks (e.g., 'instrument_seg_composite').")

        self.root_dir = "C:/Users/dsumm/OneDrive/Documents/UMD ENPM Robotics_
↳Files/BIOE658B (Intro to Medical Image Analysis)/Project/dataset/train/"
        self.label_subdir = label_subdir

        # Recursively walk through directory to find left frame image paths and_
↳GT image paths
        for subdir, dirs, files in os.walk(self.root_dir):
            if 'left_frames' in subdir:
                #print("Hit!")
                for file in sorted(files):
                    if file.endswith(('.png', '.jpg', '.jpeg')):
                        img_path = os.path.join(subdir, file)
                        #print(img_path)

                        gt_root = subdir.replace('left_frames', 'ground_truth')
                        mask_path = os.path.join(gt_root, self.label_subdir,
↳file)

                        if os.path.exists(mask_path):
                            #print("Hit!")
                            self.data.append({"image": img_path, "label":
↳mask_path}) # Dictionary for MONAI compatability

        transforms_list = [
            LoadImaged(keys=["image", "label"]), # Loads_
↳image data and metadata from file path dictionaries
            EnsureChannelFirstd(keys=["image", "label"]), #
↳Adjust or add the channel dimension of input data to ensure channel_first_
↳shape

            # Images are of nominal size 1280x1024 --> resizing for memory_
↳efficiency

```

```

        CenterSpatialCropd(keys=["image", "label"], roi_size=(1024, 1280)),
    ↪ # Cropping background padding from images
        Resized(keys=["image", "label"], spatial_size=(256, 320)),
    ↪ # Imported images are of various sizes: standardize to 320,256

        # Apply data augmentation techniques
        RandFlipd(keys=["image", "label"], prob=0.3, spatial_axis=1),
    ↪ # Horizontal axis flip imposed w/ 30% prob
        #RandRotate90d(keys=["image", "label"], prob=0.3, max_k=3),
    ↪ # Random 90° rotation imposed w/ 30% prob
        RandZoomd(keys=["image", "label"], prob=0.3, min_zoom=0.75,
    ↪ max_zoom=1.25), # Zoom range (+/-25%) imposed w/ 30% prob
        #RandAdjustContrastd(keys=["image"], prob=0.3, gamma=(0.75, 1.25)),
    ↪ # Contrast variation (+/-25%) imposed w/ 30% prob

        ScaleIntensityd(keys=["image"]), # Scale
    ↪ the intensity of input image to the value range 0-1
        ToTensord(keys=["image", "label"]), #
    ↪ Ensure data is of tensor type for pytorch usage
    ]

    # Additional conditional transforms based on label_subdir
    if label_subdir == "binary_composite":
        transforms_list.append(AsDiscreted(keys=["label"], threshold=0.5))
    ↪ # Binary threshold for binary seg
    elif label_subdir == "part_seg_composite":
        transforms_list.append(AsDiscreted(keys=["label"], to_onehot=5))
    ↪ # 5 individual class labels for instrument independent part seg
    elif label_subdir == "instrument_seg_composite":
        transforms_list.append(AsDiscreted(keys=["label"], to_onehot=8))
    ↪ # 8 individual class labels for part independent instrument seg
    elif label_subdir == "instrument_part_seg_composite":
        transforms_list.append(AsDiscreted(keys=["label"], to_onehot=21))
    ↪ # 26 individual class labels for instrument & part seg

    # Imposing MONAI transforms
    # https://docs.monai.io/en/stable/transforms.html
    self.transform = Compose(transforms_list)

    def __len__(self):
        # Returns number of imported samples
        length = len(self.data)
        return length

    def __getitem__(self, idx):
        # Return transformed sample from the dataset as dictated by the index

```

```

sample = self.data[idx]
return self.transform(sample)

```

```

[20]: class MONAIDataLoader(LightningDataModule):
    def __init__(self, dataset=None, batch_size: int = None, img_size: int = 1,
    ↪None, dimensions:int = None):
        super().__init__()
        if dataset is None:
            raise ValueError("No dataset given!")
        else:
            self.dataset = dataset

            self.train, self.val = random_split(self.dataset, [int(len(self.
    ↪dataset) * 0.8), len(self.dataset) - int(len(self.dataset) * 0.8)])
            self.batch_size = batch_size
            #self.num_workers = 2
            self.pin_memory = True
            #self.persistent_workers = True
            print(f"Train dataset size: {len(self.train)}")
            print(f"Validation dataset size: {len(self.val)}")

    def setup(self, stage=None):
        # required by PyTorch Lightning
        pass

    def train_dataloader(self):
        return DataLoader(self.train, batch_size=self.batch_size,
    ↪pin_memory=self.pin_memory)

    def val_dataloader(self):
        return DataLoader(self.val, batch_size=self.batch_size, pin_memory=self.
    ↪pin_memory)

    #def predict_dataloader(self):
    #    return DataLoader(self.test, batch_size=self.batch_size,
    ↪num_workers=16)

```

```

[21]: class SwinUNETR_Train(LightningModule):
    def __init__(self, img_size=(1, 3, 256, 320), batch_size=1, lr=0.001,
    ↪num_classes=1):
        super().__init__()

        self.save_hyperparameters()
        self.num_classes = num_classes
        print("num_classes", self.num_classes, num_classes, self.hparams.
    ↪num_classes)
        self.example_input_array = [torch.zeros(self.hparams.img_size)]

```

```

        self.dice_metric = DiceMetric(include_background=True,
↪reduction="mean", ignore_empty=True)
        self.iou_metric = MeanIoU(include_background=True, reduction="mean",
↪ignore_empty=True)

        # Metric tracking
        self.dice_scores = []
        self.iou_scores = []

        # Define SwinUNETR model from MONAI
        self.model = SwinUNETR(
            img_size=(256,320),
            in_channels=3,
            out_channels=self.num_classes,
            feature_size=48,                                # common starting point;
↪can increase to 96/128
            drop_rate=0.1,                                  # 10% dropout probability
            use_checkpoint=True,                             # Enable gradient
↪checkpointing to save memory
            spatial_dims = 2,
        )

        # Using combined DICE and CE loss as loss function
        # Conditional loss function based on the number of classes
        if num_classes == 1:
            self.DICE_CE_Loss = DiceCELoss(
                include_background=False, # Exclude background class
                sigmoid=True, # Use softmax for multiclass segmentation
                softmax=False, # Apply softmax for multiclass
                lambda_dice=1.0, # Adjust the weight for Dice loss
                lambda_ce=1.0, # Adjust the weight for Cross-Entropy loss
                reduction='mean' # Use mean reduction
            )
        else:
            self.DICE_CE_Loss = DiceCELoss(
                include_background=False, # Exclude background class
                sigmoid=False, # Use softmax for multiclass segmentation
                softmax=True, # Apply softmax for multiclass
                lambda_dice=1.0, # Adjust the weight for Dice loss
                lambda_ce=1.0, # Adjust the weight for Cross-Entropy loss
                reduction='mean' # Use mean reduction
            )

        # Tracking losses for matplotlib
        self.train_losses = []
        self.val_losses = []

```

```

        # For storing images for the last epoch
        self.last_image = []
        self.last_pred = []
        self.last_mask = []
        self.logged_epochs = []

    # Passes model inputs through U-net to get output predictions
    def forward(self, inputs):
        outputs = self.model(inputs)
        return outputs

    def training_step(self, batch, batch_idx):
        # Gets labels for input and corresponding ground truth
        inputs, gt_input = self._prepare_batch(batch)

        # Call forward pass
        outputs = self.forward(inputs)

        # Compute DICE & CE loss based on current params
        loss = self.DICE_CE_Loss(outputs, gt_input)

        # Log DICE loss with PyTorch Lightning logger
        self.log(f"Train_Dice_CE_loss", loss, on_epoch=True, prog_bar=True)

        # Append train loss at the end of each epoch
        if batch_idx == len(batch) - 1:
            self.train_losses.append(loss.item())

        return loss

    def validation_step(self, batch, batch_idx):

        # Gets labels for input and corresponding ground truth
        inputs, gt_input = self._prepare_batch(batch)
        outputs = self.forward(inputs)
        loss = self.DICE_CE_Loss(outputs, gt_input)
        self.log("val_loss", loss, on_step=False, on_epoch=True, prog_bar=True)

        if self.hparams.num_classes == 1:
            probs = torch.sigmoid(outputs)
            preds = (probs > 0.5).float()
            # Ensure ground truth is binary (i.e., 0 or 1)
            gt_input = (gt_input > 0.5).float() # Threshold the ground truth
↪ if needed

            intersection = (preds * gt_input).sum()

```

```

        union = preds.sum() + gt_input.sum()
        bin_dice_score = 2.0 * intersection / (union + 1e-8) # Avoid
↪division by zero
        # IoU score calculation for binary segmentation
        bin_iou_score = intersection / (union - intersection + 1e-8) #
↪Avoid division by zero

        self.log("val_dice", bin_dice_score, on_step=False, on_epoch=True,
↪prog_bar=True)
        self.log("val_iou", bin_iou_score, on_step=False, on_epoch=True,
↪prog_bar=True)

    else:
        probs = torch.softmax(outputs, dim=1)
        preds = torch.nn.functional.one_hot(torch.argmax(probs, dim=1),
↪num_classes=self.num_classes)
        preds = preds.permute(0, 3, 1, 2).float() # Shape: [B, C, H, W]

        self.dice_metric(y_pred=preds, y=gt_input)
        self.iou_metric(y_pred=preds, y=gt_input)

    if self.trainer.sanity_checking:
        return # skip logging during sanity check

    # Append validation loss at the end of each epoch
    if batch_idx == len(batch) - 1:
        self.val_losses.append(loss.item())

    # For binary segmentation: apply sigmoid and threshold
    if self.hparams.num_classes == 1:
        outputs = torch.sigmoid(outputs)
        outputs = (outputs > 0.5).float() # Convert probabilities to
↪binary mask

        self.dice_scores.append(bin_dice_score)
        self.iou_scores.append(bin_iou_score)

    # For multiclass segmentation: apply softmax
    else:
        outputs = torch.softmax(outputs, dim=1) # Apply softmax for
↪multi-class outputs
        dice = self.dice_metric.aggregate()[0].item()
        #print("Dice", dice)
        iou = self.iou_metric.aggregate()[0].item()
        #print("IOU", iou)
        self.dice_metric.reset()
        self.iou_metric.reset()

```

```

        self.dice_scores.append(dice)
        self.iou_scores.append(iou)
        self.log("val_dice", dice, on_step=False, on_epoch=True,
↪prog_bar=True)
        self.log("val_iou", iou, on_step=False, on_epoch=True,
↪prog_bar=True)

        # Normalize and convert tensor to 3 channels (RGB) for visualization
        def process(last):
            # Detach from cpu to not interrupt training
            # https://stackoverflow.com/questions/63582590/
↪why-do-we-call-detach-before-calling-numpy-on-a-pytorch-tensor
            last = last[0].detach().cpu()

            # Min max normalization
            # https://www.codecademy.com/article/normalization
            last = (last - last.min()) / (last.max() - last.min() + 1e-8)

            # If grayscale, reshape last image to RGB for display by
↪replicating gray value twice
            # https://discuss.pytorch.org/t/convert-grayscale-images-to-rgb/
↪113422
            return last.repeat(3, 1, 1) if last.shape[0] == 1 else last

        current_epoch = self.current_epoch
        total_epochs = self.trainer.max_epochs
        #print("TE", total_epochs)

        if current_epoch == 0 or current_epoch == total_epochs - 1 or
↪current_epoch == total_epochs // 2:
            self.last_image.append(process(inputs))
            self.last_pred.append(process(outputs))
            self.last_mask.append(process(gt_input))
            self.logged_epochs.append(current_epoch)
            print(f"Logged image from epoch {current_epoch}")

        return loss

    #def predict_step(self, batch, batch_idx, dataloader_idx=0):
    #    return self(batch['image'])

    def configure_optimizers(self):
        #set optimizer
        optimizer = torch.optim.AdamW(self.parameters(), lr=self.hparams.lr,
↪weight_decay=1e-4)

```



```

        scheduler = StepLR(optimizer, step_size=5, gamma=0.5) # halve LR every
↪5 epochs
    return {
        'optimizer': optimizer,
        'lr_scheduler': {
            'scheduler': scheduler,
            'interval': 'epoch',
            'frequency': 1
        }
    }

def _prepare_batch(self, batch):
    return batch['image'], batch['label']

# Plot training and val losses when needed
def plot_losses(self):
    min_len = min(len(self.train_losses), len(self.val_losses))
    epochs = range(1, min_len + 1)

    # Plotting training vs validation loss
    plt.figure(figsize=(10, 6))
    plt.plot(epochs, self.train_losses[:len(epochs)], label="Training
↪Loss", color='blue')
    plt.plot(epochs, self.val_losses[:len(epochs)], label="Validation
↪Loss", color='orange')
    plt.title("Training vs Validation Loss")
    plt.xlabel("Epochs")
    plt.ylabel("Loss")
    plt.legend()
    plt.show()

def plot_metrics(self):
    epochs = range(1, len(self.dice_scores) + 1)

    # Convert to CPU floats if necessary
    dice = [d.cpu().item() if torch.is_tensor(d) else d for d in self.
↪dice_scores]
    iou = [i.cpu().item() if torch.is_tensor(i) else i for i in self.
↪iou_scores]

    plt.figure(figsize=(10, 6))
    plt.plot(epochs, dice, label='Dice Coefficient')
    plt.plot(epochs, iou, label='IoU')
    plt.xlabel("Epochs")
    plt.ylabel("Score")
    plt.title("Validation Metrics Over Time")
    plt.legend()

```

```

plt.show()

def plot_result_by_epoch(self):
    total_epochs = len(self.last_image)
    print("Total Epochs:", total_epochs)

    if total_epochs < 5:
        print(f"Only {total_epochs} epochs recorded, plotting all.")
        selected_epochs = list(range(total_epochs))
    else:
        print(f"{total_epochs} epochs recorded, bug in code.")

    for epoch_idx in selected_epochs:
        epoch_num = self.logged_epochs[epoch_idx] if hasattr(self,
↪ "logged_epochs") else epoch_idx
        img = self.last_image[epoch_idx]
        pred = self.last_pred[epoch_idx]
        mask = self.last_mask[epoch_idx]

        fig, ax = plt.subplots(1, 3, figsize=(12, 4))

        ax[0].imshow(np.transpose(img.numpy(), (1, 2, 0)))
        ax[0].set_title(f"Epoch {epoch_num} - Image")
        ax[0].axis("off")

        if self.hparams.num_classes == 1:
            ax[1].imshow(np.transpose(pred.numpy(), (1, 2, 0)))
            ax[1].set_title(f"Epoch {epoch_num} - Prediction")
            ax[1].axis("off")

            ax[2].imshow(np.transpose(mask.numpy(), (1, 2, 0)))
            ax[2].set_title(f"Epoch {epoch_num} - Ground Truth")
            ax[2].axis("off")
        else:
            # Define the colormap and normalization
            num_classes = self.hparams.num_classes
            cmap = plt.get_cmap('viridis', num_classes)
            bounds = np.arange(num_classes + 1) - 0.5
            norm = plt.matplotlib.colors.BoundaryNorm(bounds, cmap.N)

            # Convert one-hot encoded predictions and masks to
↪ single-channel class labels
            pred_mask = torch.argmax(pred, dim=0).cpu().numpy()
            true_mask = torch.argmax(mask, dim=0).cpu().numpy()

            # Apply consistent colormap and normalization
            im1 = ax[1].imshow(pred_mask, cmap=cmap, norm=norm)

```

```

ax[1].set_title(f"Epoch {epoch_num} - Prediction")
ax[1].axis("off")

im2 = ax[2].imshow(true_mask, cmap=cmap, norm=norm)
ax[2].set_title(f"Epoch {epoch_num} - Ground Truth")
ax[2].axis("off")

im_for_cbar = im1 # just need one mappable

# Adjust layout to leave space at the bottom
fig.subplots_adjust(bottom=0.25) # tweak this if labels get cut
↳off

# Add a new axis below the plots for the colorbar
cbar_ax = fig.add_axes([0.1, 0.1, 0.8, 0.10]) # [left, bottom,
↳width, height]
cbar = fig.colorbar(im_for_cbar, cax=cbar_ax,
↳orientation='horizontal', ticks=np.arange(num_classes))

# Set class labels
if num_classes == 5:
    cbar.ax.set_xticklabels(['Background', 'Shaft', 'Wrist',
↳'Claspers', 'Probe'])
elif num_classes == 8:
    cbar.ax.set_xticklabels(['Background', 'Bipolar Forceps',
↳'Prograsp Forceps', 'Large Needle Driver',
↳'Vessel Sealer', 'Grasping
↳Retractor', 'Monopolar Curved Scissors', 'Other'])

plt.setp(cbar.ax.get_xticklabels(), rotation=30,
↳ha="right", rotation_mode="anchor")
elif num_classes == 21:
    cbar.ax.set_xticklabels([
        "Background",
        "Bipolar Forceps Shaft", "Bipolar Forceps Wrist",
↳"Bipolar Forceps Claspers",
        "Prograsp Forceps Shaft", "Prograsp Forceps Wrist",
↳"Prograsp Forceps Claspers",
        "Large Needle Driver Shaft", "Large Needle Driver
↳Wrist", "Large Needle Driver Claspers",
        "Vessel Sealer Shaft", "Vessel Sealer Wrist", "Vessel
↳Sealer Claspers",
        "Grasping Retractor Shaft", "Grasping Retractor Wrist",
↳"Grasping Retractor Claspers",
        "Monopolar Curved Scissors Shaft", "Monopolar Curved
↳Scissors Wrist", "Monopolar Curved Scissors Claspers",

```

```

        "Other Probe", "Other Probe"
    ])
    plt.setp(cbar.ax.get_xticklabels(), rotation=45,
    ↪ha="right", rotation_mode="anchor")

    cbar.set_label('Class ID')

    plt.show()

```

```

[22]: # Generate datasets, loaders, and models for basic UNet
binary_endo_images = EndoVis2017Dataset(label_subdir='binary_composite')
part_seg_endo_images = EndoVis2017Dataset(label_subdir='part_seg_composite')
instr_seg_endo_images =
    ↪EndoVis2017Dataset(label_subdir='instrument_seg_composite')
part_instr_seg_endo_images =
    ↪EndoVis2017Dataset(label_subdir='instrument_part_seg_composite')

binary_endo_data = MONAIDataloader(dataset=binary_endo_images, batch_size=10)
    ↪# batch size should be divisible, ie. 50 images and bs 20 wont work
part_seg_endo_data = MONAIDataloader(dataset=part_seg_endo_images,
    ↪batch_size=10)
instr_seg_endo_data = MONAIDataloader(dataset=instr_seg_endo_images,
    ↪batch_size=5)
part_instr_seg_endo_data = MONAIDataloader(dataset=part_instr_seg_endo_images,
    ↪batch_size=5)

binary_SwinUNETR_model = SwinUNETR_Train(num_classes=1)
part_seg_SwinUNETR_model = SwinUNETR_Train(num_classes=5)
instr_seg_SwinUNETR_model = SwinUNETR_Train(num_classes=8)
part_instr_seg_SwinUNETR_model = SwinUNETR_Train(num_classes=21)

```

```

Train dataset size: 1440
Validation dataset size: 360
Train dataset size: 1440
Validation dataset size: 360
Train dataset size: 1440
Validation dataset size: 360
Train dataset size: 1440
Validation dataset size: 360
num_classes 1 1 1
num_classes 5 5 5
num_classes 8 8 8
num_classes 21 21 21

```

```

[23]: if __name__ == "__main__":

    logger = TensorBoardLogger("tb_logs", name="swinunetr_binary_seg")

```

```

early_stop_callback = EarlyStopping(
    monitor="Train_Dice_CE_loss",          # metric name from self.log
    mode="min",                            # because lower loss is better
    patience=5,                            # epochs to wait before stopping
    verbose=True
)

checkpoint_callback = ModelCheckpoint(
    monitor="Train_Dice_CE_loss",
    mode="min",
    save_top_k=1,
    dirpath="checkpoints/",
    filename="best-binary-seg-SwinUNETR",
)

trainer = Trainer(
    precision="16-mixed",
    accelerator="gpu",
    max_epochs=20,
    #limit_train_batches=0.1, # or 0.1 to use 10%
    logger=logger,
    callbacks=[early_stop_callback, checkpoint_callback],
)

start_train = time.time()
trainer.fit(
    model=binary_SwinUNETR_model,
    datamodule=binary_endo_data
)
end_train = time.time()
print(f"Training time: {(end_train - start_train)/60:.2f} minutes")

# Plot the overlaid training and val loss curves per epoch
binary_SwinUNETR_model.plot_losses()

# Plot the IOU and DSC curves per epoch
binary_SwinUNETR_model.plot_metrics()

# Plot images from last epoch
binary_SwinUNETR_model.plot_result_by_epoch()

os.makedirs('SwinUNETRmodels', exist_ok=True)
# Define file names with paths
binary_SwinUNETR_model_filename = 'SwinUNETRmodels/binary_SwinUNETR_model.
pth'
# Save the model parameters

```

```

    torch.save(binary_SwinUNETR_model.state_dict(),
↪binary_SwinUNETR_model_filename)
    print("Model saved in the 'SwinUNETRmodels' directory!")

```

Using 16bit Automatic Mixed Precision (AMP)

GPU available: True (cuda), used: True

TPU available: False, using: 0 TPU cores

HPU available: False, using: 0 HPUs

LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]

	Name	Type	Params	Mode	In sizes	Out sizes
0	model	SwinUNETR	25.1 M	train	[1, 3, 256, 320]	[1, 1, 256, 320]
1	DICE_CE_Loss	DiceCELoss	0	train	?	?

```

25.1 M    Trainable params
0         Non-trainable params
25.1 M    Total params
100.557   Total estimated model params size (MB)
274       Modules in train mode
0         Modules in eval mode

```

Sanity Checking: | | 0/? [00:00<?, ?it/s]

Training: | | 0/? [00:00<?, ?it/s]

Validation: | | 0/? [00:00<?, ?it/s]

Logged image from epoch 0

Metric Train_Dice_CE_loss improved. New best score: 0.473

Validation: | | 0/? [00:00<?, ?it/s]

Metric Train_Dice_CE_loss improved by 0.224 >= min_delta = 0.0. New best score: 0.249

Validation: | | 0/? [00:00<?, ?it/s]

Metric Train_Dice_CE_loss improved by 0.059 >= min_delta = 0.0. New best score: 0.190

Validation: | | 0/? [00:00<?, ?it/s]

Metric Train_Dice_CE_loss improved by 0.021 >= min_delta = 0.0. New best score: 0.169

Validation: | | 0/? [00:00<?, ?it/s]

Metric Train_Dice_CE_loss improved by 0.012 >= min_delta = 0.0. New best score: 0.158

Validation: | | 0/? [00:00<?, ?it/s]
Metric Train_Dice_CE_loss improved by 0.023 >= min_delta = 0.0. New best score: 0.135

Validation: | | 0/? [00:00<?, ?it/s]
Metric Train_Dice_CE_loss improved by 0.007 >= min_delta = 0.0. New best score: 0.128

Validation: | | 0/? [00:00<?, ?it/s]
Metric Train_Dice_CE_loss improved by 0.007 >= min_delta = 0.0. New best score: 0.121

Validation: | | 0/? [00:00<?, ?it/s]
Metric Train_Dice_CE_loss improved by 0.005 >= min_delta = 0.0. New best score: 0.117

Validation: | | 0/? [00:00<?, ?it/s]
Metric Train_Dice_CE_loss improved by 0.001 >= min_delta = 0.0. New best score: 0.116

Validation: | | 0/? [00:00<?, ?it/s]
Logged image from epoch 10

Metric Train_Dice_CE_loss improved by 0.010 >= min_delta = 0.0. New best score: 0.106

Validation: | | 0/? [00:00<?, ?it/s]
Metric Train_Dice_CE_loss improved by 0.002 >= min_delta = 0.0. New best score: 0.103

Validation: | | 0/? [00:00<?, ?it/s]
Metric Train_Dice_CE_loss improved by 0.000 >= min_delta = 0.0. New best score: 0.103

Validation: | | 0/? [00:00<?, ?it/s]
Metric Train_Dice_CE_loss improved by 0.001 >= min_delta = 0.0. New best score: 0.102

Validation: | | 0/? [00:00<?, ?it/s]
Metric Train_Dice_CE_loss improved by 0.001 >= min_delta = 0.0. New best score: 0.101

Validation: | | 0/? [00:00<?, ?it/s]
Metric Train_Dice_CE_loss improved by 0.005 >= min_delta = 0.0. New best score: 0.096

Validation: | | 0/? [00:00<?, ?it/s]

Metric Train_Dice_CE_loss improved by 0.003 >= min_delta = 0.0. New best score: 0.093

Validation: | | 0/? [00:00<?, ?it/s]

Metric Train_Dice_CE_loss improved by 0.001 >= min_delta = 0.0. New best score: 0.092


Validation: | | 0/? [00:00<?, ?it/s]


Validation: | | 0/? [00:00<?, ?it/s]

Logged image from epoch 19

`Trainer.fit` stopped: `max_epochs=20` reached.

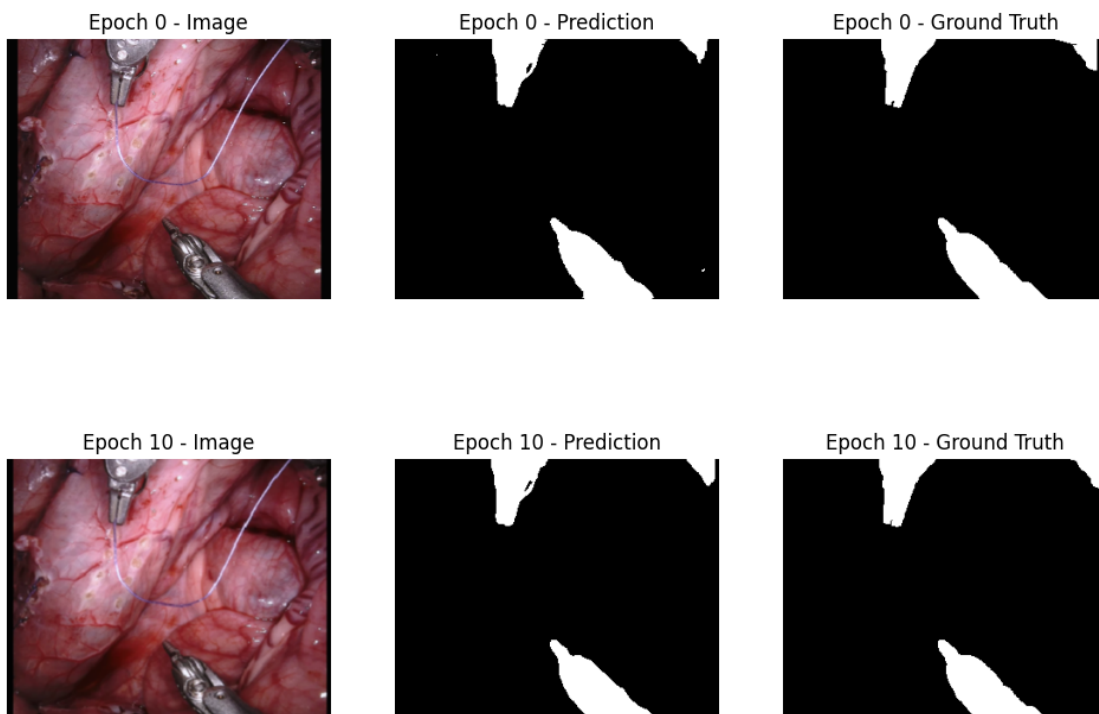
Training time: 42.99 minutes






Total Epochs: 3

Only 3 epochs recorded, plotting all.





Model saved in the 'SwinUNETRmodels' directory!

```
[24]: if __name__ == "__main__":

    logger = TensorBoardLogger("tb_logs", name="swinunetr_part_seg")

    early_stop_callback = EarlyStopping(
        monitor="Train_Dice_CE_loss",          # metric name from self.log
        mode="min",                            # because lower loss is better
        patience=5,                            # epochs to wait before stopping
        verbose=True
    )

    checkpoint_callback = ModelCheckpoint(
        monitor="Train_Dice_CE_loss",
        mode="min",
        save_top_k=1,
        dirpath="checkpoints/",
        filename="best-part-seg-SwinUNETR",
    )

    trainer = Trainer(
        precision="16-mixed",
        accelerator="gpu",
        max_epochs=20,
        #limit_train_batches=0.1, # or 0.1 to use 10%
        logger=logger,
        callbacks=[early_stop_callback, checkpoint_callback],
    )

    start_train = time.time()
    trainer.fit(
        model=part_seg_SwinUNETR_model,
        datamodule=part_seg_endo_data
```

```

)
end_train = time.time()
print(f"Training time: {(end_train - start_train)/60:.2f} minutes")

# Plot the overlaid training and val loss curves per epoch
part_seg_SwinUNETR_model.plot_losses()

# Plot the IOU and DSC curves per epoch
part_seg_SwinUNETR_model.plot_metrics()

# Plot images from last epoch
part_seg_SwinUNETR_model.plot_result_by_epoch()

os.makedirs('SwinUNETRmodels', exist_ok=True)
# Define file names with paths
part_seg_SwinUNETR_model_filename = 'SwinUNETRmodels/
↪part_seg_SwinUNETR_model.pth'
# Save the model parameters
torch.save(part_seg_SwinUNETR_model.state_dict(), ↪
↪part_seg_SwinUNETR_model_filename)
print("Model saved in the 'SwinUNETRmodels' directory!")

```

Using 16bit Automatic Mixed Precision (AMP)

GPU available: True (cuda), used: True

TPU available: False, using: 0 TPU cores

HPU available: False, using: 0 HPUs

LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]

	Name	Type	Params	Mode	In sizes	Out sizes
0	model	SwinUNETR	25.1 M	train	[1, 3, 256, 320]	[1, 5, 256, 320]
1	DICE_CE_Loss	DiceCELoss	0	train	?	?

```

-----
25.1 M   Trainable params
0        Non-trainable params
25.1 M   Total params
100.558  Total estimated model params size (MB)
274      Modules in train mode
0        Modules in eval mode

```

Sanity Checking: | | 0/? [00:00<?, ?it/s]

Training: | | 0/? [00:00<?, ?it/s]

Validation: | | 0/? [00:00<?, ?it/s]

Logged image from epoch 0

Metric Train_Dice_CE_loss improved. New best score: 1.113

Validation: | | 0/? [00:00<?, ?it/s]

Metric Train_Dice_CE_loss improved by 0.281 >= min_delta = 0.0. New best score: 0.832

Validation: | | 0/? [00:00<?, ?it/s]

Metric Train_Dice_CE_loss improved by 0.090 >= min_delta = 0.0. New best score: 0.742

Validation: | | 0/? [00:00<?, ?it/s]

Metric Train_Dice_CE_loss improved by 0.042 >= min_delta = 0.0. New best score: 0.700

Validation: | | 0/? [00:00<?, ?it/s]

Metric Train_Dice_CE_loss improved by 0.042 >= min_delta = 0.0. New best score: 0.658

Validation: | | 0/? [00:00<?, ?it/s]

Metric Train_Dice_CE_loss improved by 0.050 >= min_delta = 0.0. New best score: 0.608

Validation: | | 0/? [00:00<?, ?it/s]

Metric Train_Dice_CE_loss improved by 0.022 >= min_delta = 0.0. New best score: 0.586

Validation: | | 0/? [00:00<?, ?it/s]

Metric Train_Dice_CE_loss improved by 0.011 >= min_delta = 0.0. New best score: 0.575

Validation: | | 0/? [00:00<?, ?it/s]

Metric Train_Dice_CE_loss improved by 0.009 >= min_delta = 0.0. New best score: 0.565

Validation: | | 0/? [00:00<?, ?it/s]

Metric Train_Dice_CE_loss improved by 0.014 >= min_delta = 0.0. New best score: 0.551

Validation: | | 0/? [00:00<?, ?it/s]



Logged image from epoch 10

Metric Train_Dice_CE_loss improved by 0.021 >= min_delta = 0.0. New best score: 0.531

Validation: | | 0/? [00:00<?, ?it/s]

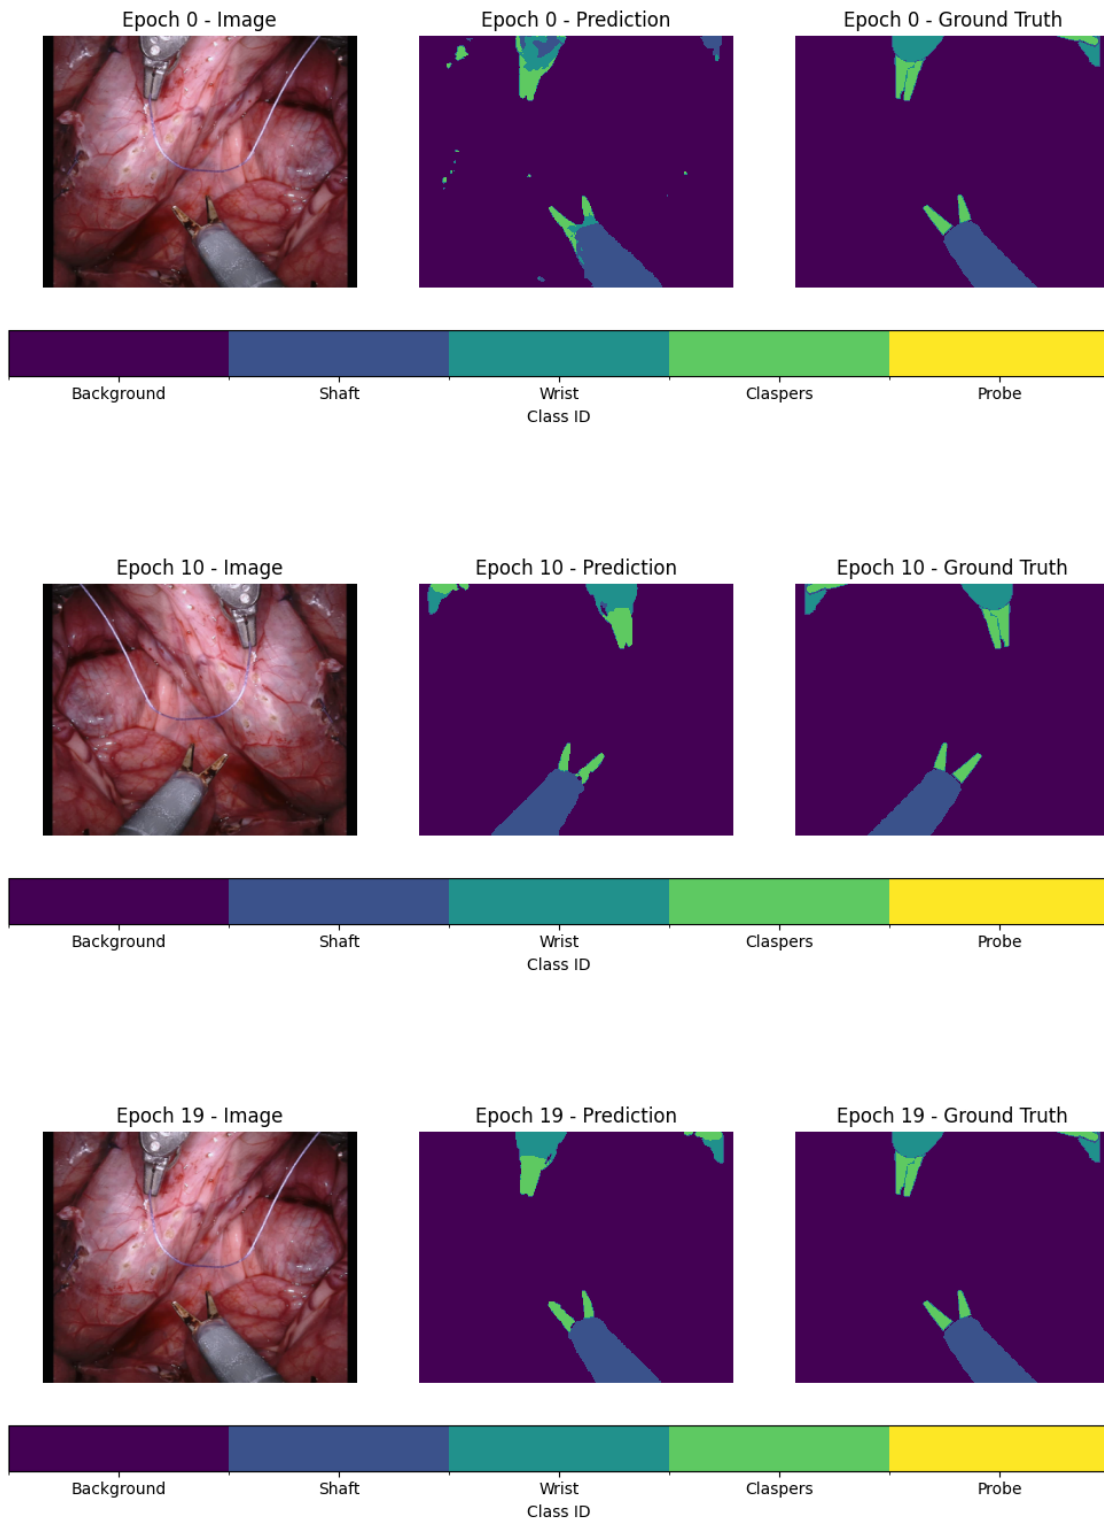
Metric Train_Dice_CE_loss improved by 0.008 >= min_delta = 0.0. New best score: 0.523

```
Validation: |           | 0/? [00:00<?, ?it/s]
Metric Train_Dice_CE_loss improved by 0.004 >= min_delta = 0.0. New best score:
0.518
Validation: |           | 0/? [00:00<?, ?it/s]
Metric Train_Dice_CE_loss improved by 0.005 >= min_delta = 0.0. New best score:
0.513
Validation: |           | 0/? [00:00<?, ?it/s]
Metric Train_Dice_CE_loss improved by 0.006 >= min_delta = 0.0. New best score:
0.507
Validation: |           | 0/? [00:00<?, ?it/s]
Metric Train_Dice_CE_loss improved by 0.014 >= min_delta = 0.0. New best score:
0.494
Validation: |           | 0/? [00:00<?, ?it/s]
Metric Train_Dice_CE_loss improved by 0.003 >= min_delta = 0.0. New best score:
0.491
Validation: |           | 0/? [00:00<?, ?it/s]
Metric Train_Dice_CE_loss improved by 0.003 >= min_delta = 0.0. New best score:
0.488
Validation: |           | 0/? [00:00<?, ?it/s]
Metric Train_Dice_CE_loss improved by 0.001 >= min_delta = 0.0. New best score:
0.487
Validation: |           | 0/? [00:00<?, ?it/s]
Logged image from epoch 19
Metric Train_Dice_CE_loss improved by 0.006 >= min_delta = 0.0. New best score:
0.481
`Trainer.fit` stopped: `max_epochs=20` reached.
Training time: 47.25 minutes
```



Total Epochs: 3

Only 3 epochs recorded, plotting all.



Model saved in the 'SwinUNETRmodels' directory!

```
[25]: if __name__ == "__main__":

    logger = TensorBoardLogger("tb_logs", name="swinunetr_instrument_seg")

    early_stop_callback = EarlyStopping(
        monitor="Train_Dice_CE_loss",          # metric name from self.log
        mode="min",                            # because lower loss is better
        patience=5,                            # epochs to wait before stopping
        verbose=True
    )

    checkpoint_callback = ModelCheckpoint(
        monitor="Train_Dice_CE_loss",
        mode="min",
        save_top_k=1,
        dirpath="checkpoints/",
        filename="best-instr-seg-SwinUNETR",
    )

    trainer = Trainer(
        precision="16-mixed",
        accelerator="gpu",
        max_epochs=20,
        #limit_train_batches=0.1, # or 0.1 to use 10%
        logger=logger,
        callbacks=[early_stop_callback, checkpoint_callback],
    )

    start_train = time.time()
    trainer.fit(
        model=instr_seg_SwinUNETR_model,
        datamodule=instr_seg_endo_data
    )
    end_train = time.time()
    print(f"Training time: {(end_train - start_train)/60:.2f} minutes")

    # Plot the overlaid training and val loss curves per epoch
    instr_seg_SwinUNETR_model.plot_losses()

    # Plot the IOU and DSC curves per epoch
    instr_seg_SwinUNETR_model.plot_metrics()

    # Plot images from last epoch
    instr_seg_SwinUNETR_model.plot_result_by_epoch()
```



```

os.makedirs('SwinUNETRmodels', exist_ok=True)
instr_seg_SwinUNETR_model_filename = 'SwinUNETRmodels/
↳instr_seg_SwinUNETR_model.pth'
torch.save(instr_seg_SwinUNETR_model.state_dict(),
↳instr_seg_SwinUNETR_model_filename)
print("Model saved in the 'SwinUNETRmodels' directory!")

```

Using 16bit Automatic Mixed Precision (AMP)

GPU available: True (cuda), used: True

TPU available: False, using: 0 TPU cores

HPU available: False, using: 0 HPUs

Checkpoint directory C:\Users\dsumm\OneDrive\Documents\UMD ENPM Robotics Files\BIOE658B (Intro to Medical Image Analysis)\Project\code\checkpoints exists and is not empty.

LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]

	Name	Type	Params	Mode	In sizes	Out sizes
0	model	SwinUNETR	25.1 M	train	[1, 3, 256, 320]	[1, 8, 256, 320]
1	DICE_CE_Loss	DiceCELoss	0	train	?	?

```

-----
25.1 M    Trainable params
0         Non-trainable params
25.1 M    Total params
100.558   Total estimated model params size (MB)
274       Modules in train mode
0         Modules in eval mode

```

Sanity Checking: | | 0/? [00:00<?, ?it/s]

The 'val_dataloader' does not have many workers which may be a bottleneck.

Consider increasing the value of the `num_workers` argument` to `num_workers=31` in the `DataLoader` to improve performance.

The 'train_dataloader' does not have many workers which may be a bottleneck.

Consider increasing the value of the `num_workers` argument` to `num_workers=31` in the `DataLoader` to improve performance.

Training: | | 0/? [00:00<?, ?it/s]

Validation: | | 0/? [00:00<?, ?it/s]

Logged image from epoch 0

Metric Train_Dice_CE_loss improved. New best score: 1.311

Validation: | | 0/? [00:00<?, ?it/s]

Metric Train_Dice_CE_loss improved by 0.186 >= min_delta = 0.0. New best score:

1.125

Validation: | | 0/? [00:00<?, ?it/s]

Metric Train_Dice_CE_loss improved by 0.093 >= min_delta = 0.0. New best score: 1.033

Validation: | | 0/? [00:00<?, ?it/s]

Metric Train_Dice_CE_loss improved by 0.064 >= min_delta = 0.0. New best score: 0.969

Validation: | | 0/? [00:00<?, ?it/s]

Metric Train_Dice_CE_loss improved by 0.044 >= min_delta = 0.0. New best score: 0.925

Validation: | | 0/? [00:00<?, ?it/s]

Metric Train_Dice_CE_loss improved by 0.044 >= min_delta = 0.0. New best score: 0.881

Validation: | | 0/? [00:00<?, ?it/s]

Metric Train_Dice_CE_loss improved by 0.011 >= min_delta = 0.0. New best score: 0.870

Validation: | | 0/? [00:00<?, ?it/s]

Metric Train_Dice_CE_loss improved by 0.011 >= min_delta = 0.0. New best score: 0.859

Validation: | | 0/? [00:00<?, ?it/s]

Metric Train_Dice_CE_loss improved by 0.005 >= min_delta = 0.0. New best score: 0.854

Validation: | | 0/? [00:00<?, ?it/s]

Metric Train_Dice_CE_loss improved by 0.005 >= min_delta = 0.0. New best score: 0.849

Validation: | | 0/? [00:00<?, ?it/s]

Logged image from epoch 10

Metric Train_Dice_CE_loss improved by 0.023 >= min_delta = 0.0. New best score: 0.826

Validation: | | 0/? [00:00<?, ?it/s]

Metric Train_Dice_CE_loss improved by 0.005 >= min_delta = 0.0. New best score: 0.822

Validation: | | 0/? [00:00<?, ?it/s]

Metric Train_Dice_CE_loss improved by 0.006 >= min_delta = 0.0. New best score: 0.815

Validation: | | 0/? [00:00<?, ?it/s]

Metric Train_Dice_CE_loss improved by 0.004 >= min_delta = 0.0. New best score: 0.811

Validation: | | 0/? [00:00<?, ?it/s]

Metric Train_Dice_CE_loss improved by 0.006 >= min_delta = 0.0. New best score: 0.805

Validation: | | 0/? [00:00<?, ?it/s]

Metric Train_Dice_CE_loss improved by 0.008 >= min_delta = 0.0. New best score: 0.797

Validation: | | 0/? [00:00<?, ?it/s]

Metric Train_Dice_CE_loss improved by 0.003 >= min_delta = 0.0. New best score: 0.795

Validation: | | 0/? [00:00<?, ?it/s]

Metric Train_Dice_CE_loss improved by 0.002 >= min_delta = 0.0. New best score: 0.793

Validation: | | 0/? [00:00<?, ?it/s]

Metric Train_Dice_CE_loss improved by 0.003 >= min_delta = 0.0. New best score: 0.790



Validation: | | 0/? [00:00<?, ?it/s]

Logged image from epoch 19

Metric Train_Dice_CE_loss improved by 0.001 >= min_delta = 0.0. New best score: 0.789

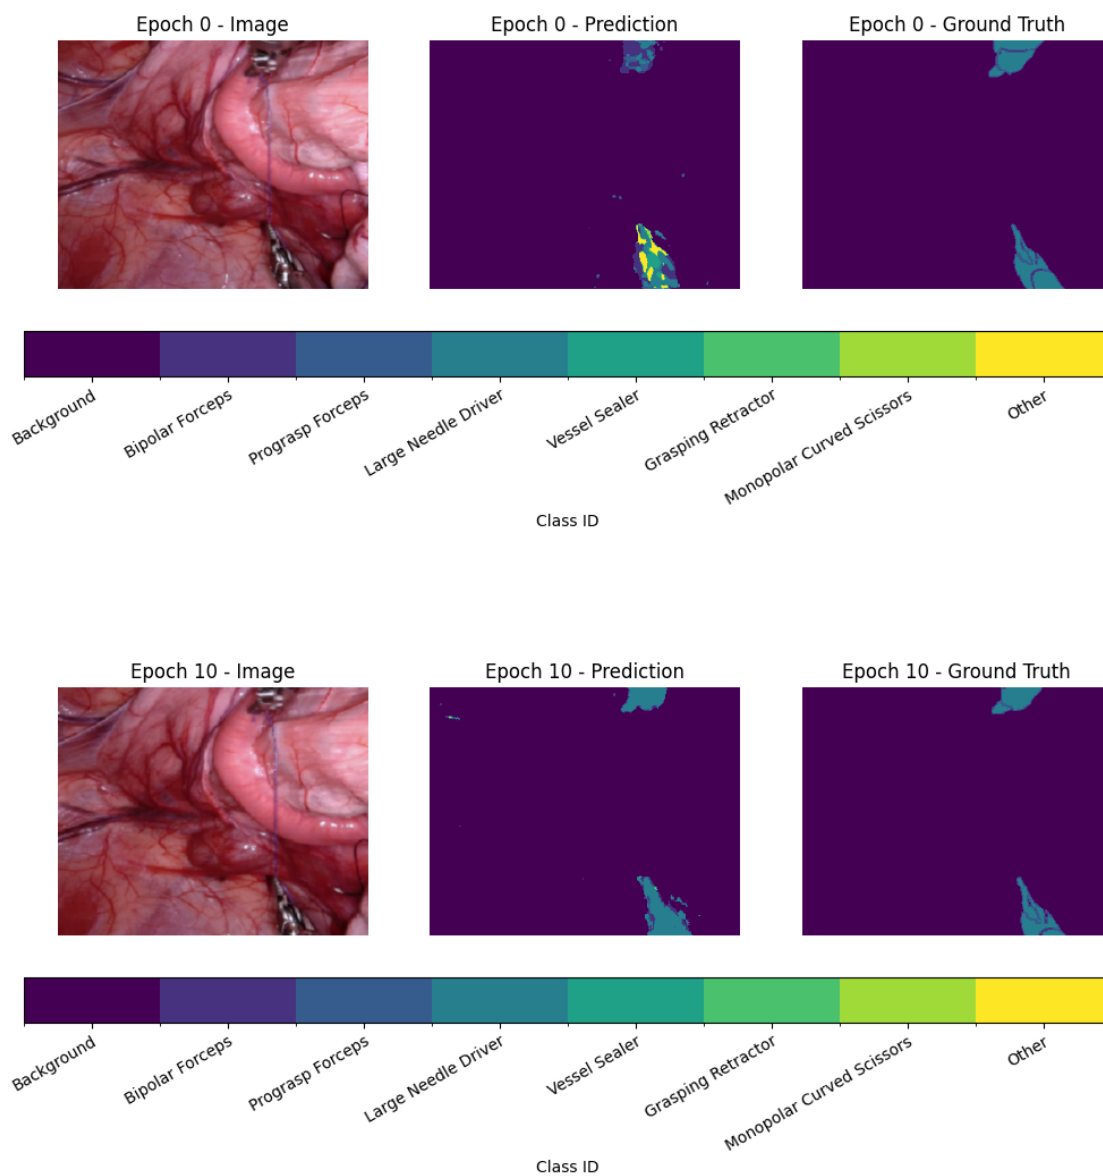
`Trainer.fit` stopped: `max_epochs=20` reached.

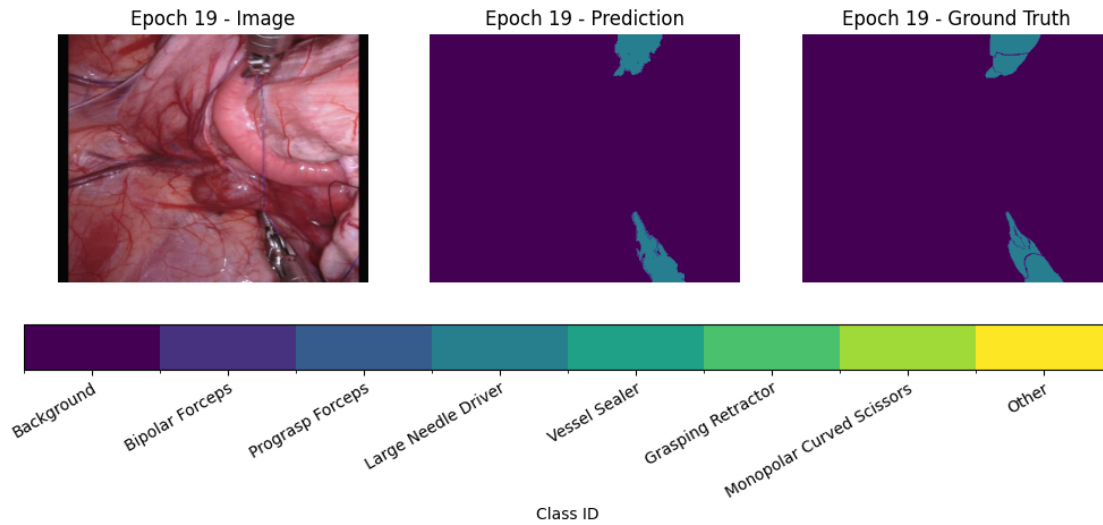
Training time: 45.86 minutes



Total Epochs: 3

Only 3 epochs recorded, plotting all.





Model saved in the 'SwinUNETRmodels' directory!

```
[26]: if __name__ == "__main__":

    logger = TensorBoardLogger("tb_logs", name="swinunetr_part_instrument_seg")

    early_stop_callback = EarlyStopping(
        monitor="Train_Dice_CE_loss",          # metric name from self.log
        mode="min",                            # because lower loss is better
        patience=5,                            # epochs to wait before stopping
        verbose=True
    )

    checkpoint_callback = ModelCheckpoint(
        monitor="Train_Dice_CE_loss",
        mode="min",
        save_top_k=1,
        dirpath="checkpoints/",
        filename="best-part-instr-seg-SwinUNETR",
    )

    trainer = Trainer(
        precision="16-mixed",
        accelerator="gpu",
        max_epochs=20,
        #limit_train_batches=0.1, # or 0.1 to use 10%
        logger=logger,
        callbacks=[early_stop_callback, checkpoint_callback],
    )
```

```

start_train = time.time()
trainer.fit(
    model=part_instr_seg_SwinUNETR_model,
    datamodule=part_instr_seg_endo_data
)
end_train = time.time()
print(f"Training time: {(end_train - start_train)/60:.2f} minutes")

# Plot the overlaid training and val loss curves per epoch
part_instr_seg_SwinUNETR_model.plot_losses()

# Plot the IOU and DSC curves per epoch
part_instr_seg_SwinUNETR_model.plot_metrics()

# Plot images from last epoch
part_instr_seg_SwinUNETR_model.plot_result_by_epoch()

os.makedirs('SwinUNETRmodels', exist_ok=True)
part_instr_seg_SwinUNETR_model_filename = 'SwinUNETRmodels/
↳part_instr_seg_SwinUNETR_model.pth'
torch.save(part_instr_seg_SwinUNETR_model.state_dict(),
↳part_instr_seg_SwinUNETR_model_filename)
print("Model saved in the 'SwinUNETRmodels' directory!")

```

Using 16bit Automatic Mixed Precision (AMP)
GPU available: True (cuda), used: True
TPU available: False, using: 0 TPU cores
HPU available: False, using: 0 HPUs
Checkpoint directory C:\Users\dsumm\OneDrive\Documents\UMD ENPM Robotics
Files\BIOE658B (Intro to Medical Image Analysis)\Project\code\checkpoints exists
and is not empty.
LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]

	Name	Type	Params	Mode	In sizes	Out sizes
0	model	SwinUNETR	25.1 M	train	[1, 3, 256, 320]	[1, 21, 256, 320]
1	DICE_CE_Loss	DiceCELoss	0	train	?	?

25.1 M	Trainable params
0	Non-trainable params
25.1 M	Total params
100.561	Total estimated model params size (MB)
274	Modules in train mode

```

0          Modules in eval mode
Sanity Checking: |          | 0/? [00:00<?, ?it/s]
The 'val_dataloader' does not have many workers which may be a bottleneck.
Consider increasing the value of the `num_workers` argument` to `num_workers=31`
in the `DataLoader` to improve performance.
The 'train_dataloader' does not have many workers which may be a bottleneck.
Consider increasing the value of the `num_workers` argument` to `num_workers=31`
in the `DataLoader` to improve performance.
Training: |          | 0/? [00:00<?, ?it/s]
Validation: |          | 0/? [00:00<?, ?it/s]
Logged image from epoch 0
Metric Train_Dice_CE_loss improved. New best score: 1.593
Validation: |          | 0/? [00:00<?, ?it/s]
Metric Train_Dice_CE_loss improved by 0.279 >= min_delta = 0.0. New best score:
1.314
Validation: |          | 0/? [00:00<?, ?it/s]
Metric Train_Dice_CE_loss improved by 0.086 >= min_delta = 0.0. New best score:
1.229
Validation: |          | 0/? [00:00<?, ?it/s]
Metric Train_Dice_CE_loss improved by 0.074 >= min_delta = 0.0. New best score:
1.155
Validation: |          | 0/? [00:00<?, ?it/s]
Metric Train_Dice_CE_loss improved by 0.051 >= min_delta = 0.0. New best score:
1.103
Validation: |          | 0/? [00:00<?, ?it/s]
Metric Train_Dice_CE_loss improved by 0.049 >= min_delta = 0.0. New best score:
1.055
Validation: |          | 0/? [00:00<?, ?it/s]
Metric Train_Dice_CE_loss improved by 0.022 >= min_delta = 0.0. New best score:
1.033
Validation: |          | 0/? [00:00<?, ?it/s]
Metric Train_Dice_CE_loss improved by 0.014 >= min_delta = 0.0. New best score:
1.020
Validation: |          | 0/? [00:00<?, ?it/s]
Metric Train_Dice_CE_loss improved by 0.013 >= min_delta = 0.0. New best score:
1.006

```


Validation: | | 0/? [00:00<?, ?it/s]
Metric Train_Dice_CE_loss improved by 0.008 >= min_delta = 0.0. New best score: 0.999

Validation: | | 0/? [00:00<?, ?it/s]
Logged image from epoch 10
Metric Train_Dice_CE_loss improved by 0.020 >= min_delta = 0.0. New best score: 0.979

Validation: | | 0/? [00:00<?, ?it/s]
Metric Train_Dice_CE_loss improved by 0.008 >= min_delta = 0.0. New best score: 0.971

Validation: | | 0/? [00:00<?, ?it/s]
Metric Train_Dice_CE_loss improved by 0.004 >= min_delta = 0.0. New best score: 0.967

Validation: | | 0/? [00:00<?, ?it/s]
Metric Train_Dice_CE_loss improved by 0.005 >= min_delta = 0.0. New best score: 0.961

Validation: | | 0/? [00:00<?, ?it/s]
Metric Train_Dice_CE_loss improved by 0.004 >= min_delta = 0.0. New best score: 0.957

Validation: | | 0/? [00:00<?, ?it/s]
Metric Train_Dice_CE_loss improved by 0.007 >= min_delta = 0.0. New best score: 0.950



Validation: | | 0/? [00:00<?, ?it/s]
Metric Train_Dice_CE_loss improved by 0.003 >= min_delta = 0.0. New best score: 0.947

Validation: | | 0/? [00:00<?, ?it/s]
Metric Train_Dice_CE_loss improved by 0.003 >= min_delta = 0.0. New best score: 0.944

Validation: | | 0/? [00:00<?, ?it/s]
Metric Train_Dice_CE_loss improved by 0.002 >= min_delta = 0.0. New best score: 0.942

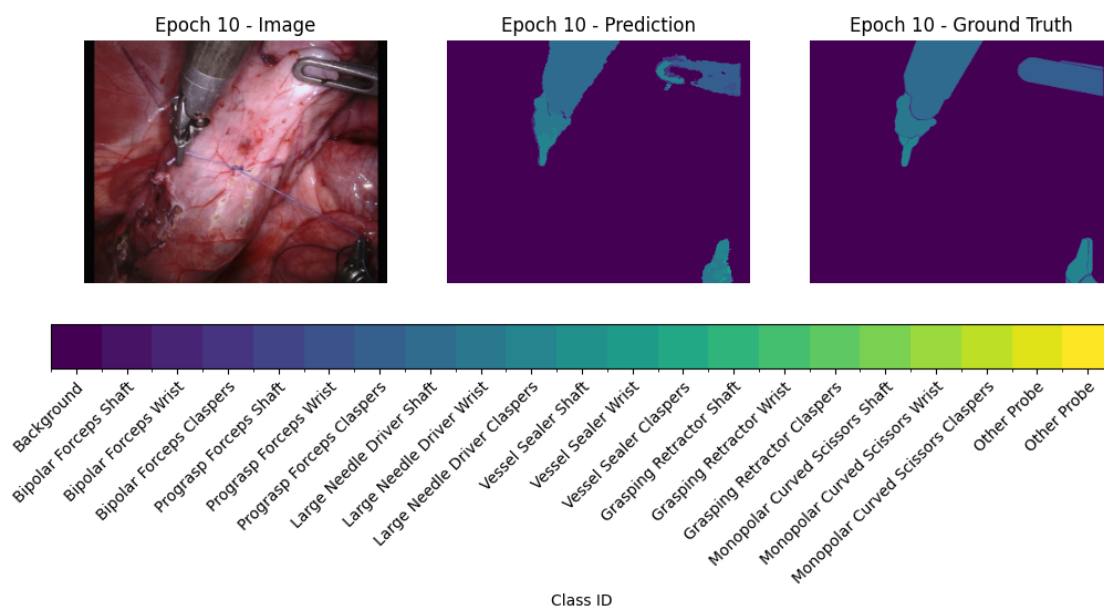
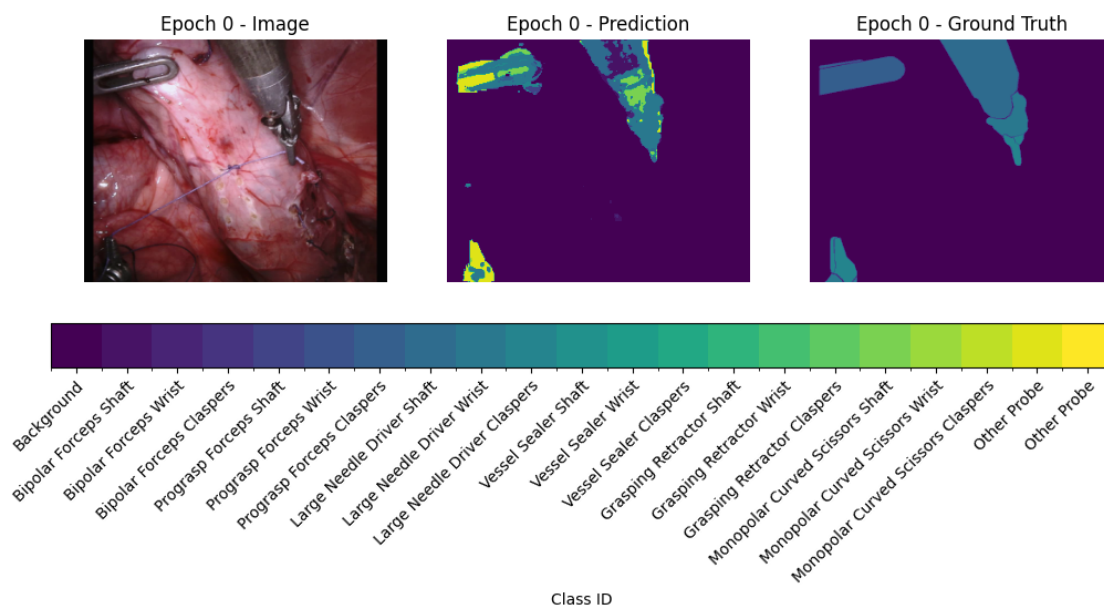
Validation: | | 0/? [00:00<?, ?it/s]
Logged image from epoch 19
Metric Train_Dice_CE_loss improved by 0.003 >= min_delta = 0.0. New best score: 0.939
`Trainer.fit` stopped: `max_epochs=20` reached.

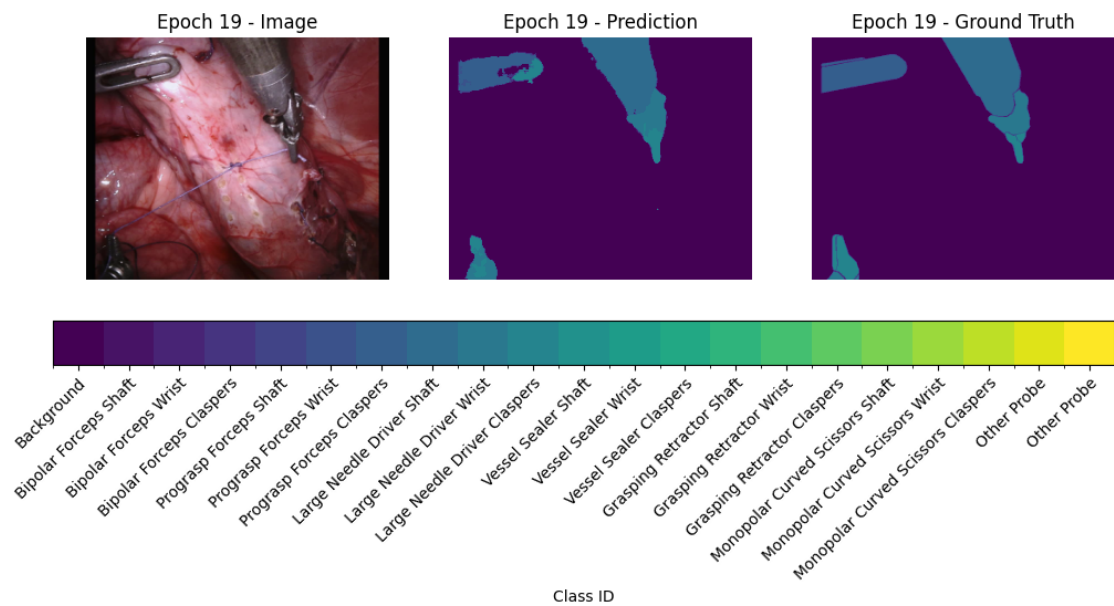
Training time: 45.99 minutes



Total Epochs: 3

Only 3 epochs recorded, plotting all.





Model saved in the 'SwinUNETRmodels' directory!